

Modelling Yearly Healthcare Costs

August 1, 2022

```
[1]: # Run this cell first!
import pandas as pd
import numpy as np
%matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
from sklearn import model_selection, linear_model, ensemble, preprocessing, \
    neural_network, metrics
```

1 Modelling Yearly Healthcare Costs

The ultimate goal of this project is to create a model to better predict yearly healthcare costs for our self-funded healthcare clients (specifically hospitals). More specifically, the model will use 36 months of claims and expense report data to predict total cost for the following year. Due to limitations in the scope of the data available, the specifications of the central question driving this analysis, and my own skills/knowledge as an amateur data scientist the model created in this project is more of a proof-of-concept than a full-fledged, deployment-ready model for predicting yearly healthcare costs. While the model is able to predict costs from historical data with sometimes striking accuracy (i.e. within +/- \$10 PEPM) it ultimately fails to predict costs with consistent accuracy—as in it is occasionally \$200+ off. That is not to say the time spent on this project was a waste; rather, this project, again, serves as proof that more advanced analytics can prove to be an incredibly valuable asset to the BUAD team. At the end of the day, previous claims experience is not enough to capture the volatility/variance of yearly healthcare costs. Given more time, I would dig deeper into the data available on Mede for our self-insured clients that could help better predict costs; but, alas, my time here has come to an end. In the following report, I'll walk through my analysis and provide a deeper, more thorough analysis/debrief in the Conclusion Section at the end of this report.

Note: Due to certain html tags utilized in this report, the html file is best viewed on a browser other than Firefox.

1.1 Introduction

For any underwriting team, calculating renewal projections can be a frustratingly difficult process. The main concern posed to underwriters is the volatility of claims experience data. Year-by-year and even month-by-month breakdowns of claims experience data yield often inconsistent numbers with frustratingly high variance. An experienced underwriter will take their industry knowledge

and apply it to historical claims experience data to mitigate the impact of volatility on their renewal projections. This project aims to leverage Machine Learning models to aid the underwriter in the renewal projection process. The model will be built using data from a variety of healthcare clients. A full list is included below: * Antelope Valley Hospital (1/1 Start | Data from 1/11 - 5/22) * Avanti/Pipeline (1/1 Start | Data from 1/14 - 12/21) * Beverly Hospital (1/1 Start | Data from 1/17 - 6/22) * CHA Hollywood Presbyterian Medical Center (1/1 Start | Data from 1/18 - 6/22) * Children's Hospital Los Angeles (7/1 Start | Data from 7/17 - 6/22) * Dameron Hospital (1/1 Start | Data from 1/11 - 6/22) * Enloe Medical Center (7/1 Start | Data from 7/12 - 6/19) * Epic Medical Group (10/1 Start | Data from 10/15 - 9/18) * Fairchild Medical Center (7/1 Start | Data from 7/13 - 6/22) * Henry Mayo Newhall Hospital (3/1 Start | Data from 3/13 - 2/22) * Huntington Memorial Hospital (1/1 Start | Data from 1/15 - 6/22) * Marshall Medical Center (11/1 Start | Data from 11/13 - 10/16) * Northern Inyo Hospital (1/1 Start | Data from 1/18 - 6/22) * Prime Healthcare Services (1/1 Start | Data from 1/11 - 6/22) * Prospect Medical Holdings (1/1 Start | Data from 1/13 - 12/20) * Salinas Valley Memorial Healthcare System (1/1 Start | Data from 1/18 - 12/21) * Tahoe Forest Health System (1/1 Start | Data from 1/15 - 6/22) * Torrance Health Association (1/1 Start | Data from 1/11 - 6/22)

The data for this project was gathered from Monthly Claims and Expense Reports (C&Es) for each of the clients enumerated above. A more detailed description of the exact data pulled will be included in the Data Section below.

Ultimately, the model will use 36 months of data to predict the next year's total healthcare cost for Keenan's self-insured healthcare clients. More specifically, it will use the actual monthly PEPM cost to predict the following year's total healthcare cost.

As an aside, there is often a misconception surrounding Data Science that the point of creating Machine Learning models is to completely strip away the human element; however the model is merely performing a series of Linear Algebra calculations that must be made sense of within the greater context of the insurance space. Moreover, it is alongside expert feedback/insight that the model was created—it is inherently human as much as it is automated. In light of that, I want to thank both Kyle for providing that expert insight throughout the process and the entire BUAD team for helping me learn more about employer-sponsored health plans and the ins-and-outs of the space—without all of you this project could not have come together.

1.1.1 The Standard Renewal Worksheet

It is worth noting that the intent of this model is not to completely replace the current processes for calculating renewal projections. Rather, it is intended to be used as a supplement with the standard projection processes to help confirm the reliability of calculated renewal projections and to help prevent volatility from blowing up a renewal projection. Some of the features included in the standard renewal worksheet are incorporated into this model—albeit not in the same way.

Trend Factor

Although not explicitly included in the model as a predictive feature, the Trend Factor (both Medical and Rx) is accounted for in the model's predictions. Fundamentally, the model identifies historical patterns and projects them into the future; as such, the Trend Factor—accounting for the increasing cost of medical treatment—is included in the model's predictions.

Administration Fees

The administration fees are included within the Adjusted Paid Claims feature used to train the model.

1.1.2 Basic Methodology

Since the goal is to (or at least attempt to) predict annual healthcare costs, the granularity of the data will be yearly (i.e. each row will represent a **Plan Year** and not a standard year). However, considering the ultimate goal is to predict healthcare costs from 36 months of previous data, the granularity of the data used to train any/all models will be 36 months (i.e. each row will represent one 36-month period).

While a more detailed description of the specific methodology will be provided within each section of the report, the basic methodology is as follows:

Exploratory Data Analysis (EDA)

Exploratory Data Analysis is the critical first-step to building Machine Learning models. By exploring the provided data to determine interesting or unusual relationships between variables or empirically confirming expected relationships between data, this stage ensures that the foundation that the model is built upon is solid. This section will focus particularly on visualizing the relationship between PEPM Cost and PMPM Cost and working to create distributions representative of certain subsets of the data that will be crucial to the process of data formation.

Modelling

Once certain relationships have been identified and/or confirmed within the data, it must be transformed and manipulated so that it is adequate for model training. Specifically, more data will need to be generated so that the sample size is sufficient. To do so, the representative distributions crafted in the previous section (representing the first four plan years for each of the clients in the data) will be sampled randomly with replacement 1000 times. A separate classification model will be used to identify and isolate points that are prone over and underestimation which will promptly be used to train two separate linear regression models for predicting yearly healthcare costs.

Note: The beginning of each section or subsection will contain a cell with all the Python user-defined functions referenced in the section. These are not to be read as text.

1.1.3 Note On Technical Specificity

Throughout this project there is an assumption that readers will have varying degrees of familiarity with the underlying statistical theory that justifies some of the approaches taken and certain programming techniques used throughout this analysis. As such, sections discussing methodology or providing justification for the use of certain statistical methods will have descriptions at varying levels of abstraction.

1.2 Exploratory Data Analysis (EDA)

```
[2]: # Python user-defined functions in alphabetical order
def calculate_weighted_pepm_avg(data, plan_year):
    '''
    Given data, data, and plan year, calculates
    the year's weighted average PEPM cost
```

```

-----
Inputs:

data (DataFrame) - DataFrame with data to calculate the year's
                    weighted average pepm cost

renewal_year (int) - Int specifying the plan year to calculate
-----
Output:

weighted_avg (float) - Calculated weighted average
'''
plan_year = data.sort_values(['Year', 'Month']).query('`Renewal Year` ==_
↳@renewal_year')
py_pepms = plan_year['PEPM'].values
py_ee_counts = plan_year['EE Count'].values
py_total_ee_count = sum(py_ee_counts)
py_props = py_ee_counts / py_total_ee_count
weighted_avg = sum(py_pepms * py_props)
return weighted_avg

def calculate_weighted_pmpm_avg(data, plan_year):
    '''
    Given data, data, and plan year, calculates
    the year's weighted average PMPM cost
    -----
    Inputs:

    data (DataFrame) - DataFrame with data to calculate the year's
                        weighted average pepm cost

    renewal_year (int) - Int specifying the plan year to calculate
    -----
    Output:

    weighted_avg (float) - Calculated weighted average
    '''
    plan_year = data.sort_values(['Year', 'Month']).query('`Renewal Year` ==_
↳@renewal_year')
    py_pepms = plan_year['PEPM'].values
    py_ee_counts = plan_year['Member Count'].values
    py_total_ee_count = sum(py_ee_counts)
    py_props = py_ee_counts / py_total_ee_count
    weighted_avg = sum(py_pepms * py_props)
    return weighted_avg

```

1.2.1 Data

As noted in the Introduction, the data used for this project is an amalgamation of data from a variety of Keenan's self-funded healthcare clients. The data will be loaded in and displayed below:

```
[3]: # Loading in the data and performing necessary data transformations
claims = pd.read_csv('Healthcare Clients.csv')
claims['PEPM'] = claims['Adjusted Total Expenses'] / claims['EE Count']
claims['PMPM'] = claims['Adjusted Total Expenses'] / claims['Member Count']
claims
```

```
[3]:
```

	Month	Year	Plan	Month	Plan	Year	Total Medical Claims \
0	1	2011		1		1	4767.00
1	2	2011		2		1	492254.80
2	3	2011		3		1	449411.00
3	4	2011		4		1	829805.58
4	5	2011		5		1	675157.46
...
1572	2	2022		2		12	700450.75
1573	3	2022		3		12	1176646.24
1574	4	2022		4		12	974612.44
1575	5	2022		5		12	954969.28
1576	6	2022		6		12	896835.62

	Total Rx Claims	Total Paid Claims	Adjusted Total Expenses	EE Count \
0	195529.47	200296.47	336582.47	1093
1	195997.93	688252.73	824538.73	1093
2	230158.09	679569.09	816977.09	1102
3	204176.23	1033981.81	1170641.81	1096
4	196067.47	871224.93	1006014.93	1081
...
1572	291079.36	991530.11	733061.77	1134
1573	241579.60	1418225.84	-69428.66	1137
1574	236503.49	1211115.93	1344295.40	1139
1575	253322.95	1208292.23	1237096.06	1144
1576	307918.72	1204754.34	1338474.19	1144

	Member Count	Client	PEPM	PMPM
0	2710	Antelope Valley	307.943705	124.200173
1	2674	Antelope Valley	754.381272	308.354050
2	2690	Antelope Valley	741.358521	303.708955
3	2683	Antelope Valley	1068.103841	436.318230
4	2613	Antelope Valley	930.633608	385.003800
...
1572	2354	Torrance Health	646.438951	311.411117
1573	2366	Torrance Health	-61.063026	-29.344320
1574	2365	Torrance Health	1180.241791	568.412431
1575	2368	Torrance Health	1081.377675	522.422323

1576 2395 Torrance Health 1169.994921 558.861875

[1577 rows x 13 columns]

This is the `claims` dataset that will ultimately be used for our model. Below is a description of each feature included in `claims`:

Month

Month of the year represented by each row with 1 corresponding to January, 2 to February, 3 to March, etc.

Year

Calendar year that corresponds to the month represented by each row of data.

Plan Month

Specifies the month number chronologically in the client's plan year. For clients with 1/1 start dates, the Plan Month feature will always be equivalent to the Month feature. As such, this feature is significant only for clients with start dates other than 1/1.

Plan Year

Specifies the plan year that corresponds to the month represented by each row. Note: the values may not correspond to the true plan year for the given client; rather, they specify the plan year for all those that are within the claims dataset.

Total Medical Claims

The total amount of Medical claims paid out in the given month. This includes Domestic Hospital Claims (both IP and OP), Non-Domestic Hospital Claims (IP and OP), and Non-Hospital Medical Claims.

Total Rx Claims

The total amount of Rx claims paid out in the given month. This includes Domestic Rx Claims and Non-Domestic Rx Claims.

Total Paid Claims

The total amount of claims (Medical and Rx) paid out in the given month or the sum of Total Medical Claims and Total Rx Claims. This does not include Stop Loss Reimbursements, Rx Rebates, or Rx Performance Guarantees.

Adjusted Total Expenses

The total amount paid out in the given month. This includes Total Paid Claims, Stop Loss Reimbursements, Rx Rebates, Rx Performance Guarantees, and Total Admin Fees. These are the Monthly Claims and Expenses values from Monthly C&Es.

EE Count

The number of employees enrolled in the client's health plan in the given month.

Member Count

The number of members enrolled in the client's health plan in the given month. The vast majority of C&Es included all of the data necessary; however, some C&Es lacked monthly member counts. In such a case, the monthly Rx Member Count was substituted as an estimate for the monthly Medical Member Count. If the monthly Rx Member Count was unavailable, a multiplier was applied to the monthly EE Count to estimate the number of enrolled members each month. These monthly multipliers were the averages of the monthly EE Count multipliers for all other months for the given client.

Client

The name of the client corresponding to the given row of data.

PEPM


The per-employee-per-month cost for the given month. This is calculated as the quotient of Adjusted Total Expenses and EE Count. Note: There were a few values < 0 which will be excluded from the dataset entirely.

PMPM

The per-member-per-month cost for the given month. This is calculated as the quotient of Adjusted Total Expenses and Member Count. Note: There were a few values < 0 which will be excluded from the dataset entirely.

1.2.2 Data Visualization

```
[4]: # Figure 1. Comparing Client PEPM Cost vs. PMPM Cost
plt.figure(figsize=(15, 10))
sns.set_style('darkgrid')
sns.scatterplot(data=claims, x='PEPM', y='PMPM', hue='Client', alpha=0.5)
plt.title('Figure 1. Comparing Client PEPM Cost vs. PMPM Cost', loc='left',
         fontdict={'fontsize': 14})
plt.xlabel('PEPM Cost ($)')
plt.ylabel('PMPM Cost ($)')
plt.xlim(0, 2000)
plt.ylim(0, 1000)
plt.show()
```





Note: For readability's sake Figure 1 limits PEPM Cost to \ \$2000 and PMPM Cost to \ \$1000.

```
[5]: plt.figure(figsize=(20, 5))

# Figure 2.1. PEPM Cost Distribution
plt.subplot(1, 2, 1)
sns.histplot(data=claims.query('PEPM > 0'), x='PEPM')
plt.xlabel('PEPM ($)')
plt.title('Figure 2.1. PEPM Cost Distribution', loc='left',
fontdict={'fontsize': 14})

# Figure 2.2. PMPM Cost Distribution
plt.subplot(1, 2, 2)
sns.histplot(data=claims.query('PMPM > 0'), x='PMPM')
plt.xlabel('PMPM ($)')
plt.title('Figure 2.2. PMPM Cost Distribution', loc='left',
fontdict={'fontsize': 14})
plt.show()
```

Figures 1, 2.1, and 2.2 all effectively demonstrate what may seem intuitive: that PEPM Cost and PMPM Cost are highly correlated. Figure 1, in essence, visualizes the EE/Member Count multiplier discussed in the Data Section above. Within each client, the points virtually form a line—indicating that the EE/Member Count multiplier does not fluctuate dramatically from month-to-month for any given client. Likewise, because the spread of the points overall is relatively narrow, it becomes apparent that the EE/Member Count multiplier does not fluctuate much between clients either. Figures 2.1 and 2.2 (taken together) communicate a similar notion. Here, both histograms display distributions that seem incredibly similar in shape. While the scales of both the x- and y-axes differ between plots, the overall shape of the distributions remains the same—suggesting the two features seem to be correlated.

```
[6]: # Figure 3. Distribution of PEPM Costs per Plan Month
plt.figure(figsize=(15, 10))
sns.boxplot(data=claims.query('PEPM > 0 and PEPM < 3000'), x='Plan Month',
            y='PEPM')
plt.yticks(np.arange(0, 3300, 100))
plt.ylabel('PEPM Cost ($)')
plt.title('Figure 3. Distribution of PEPM Costs per Plan Month', loc='left',
        fontdict={'fontsize': 14})
plt.show()
```

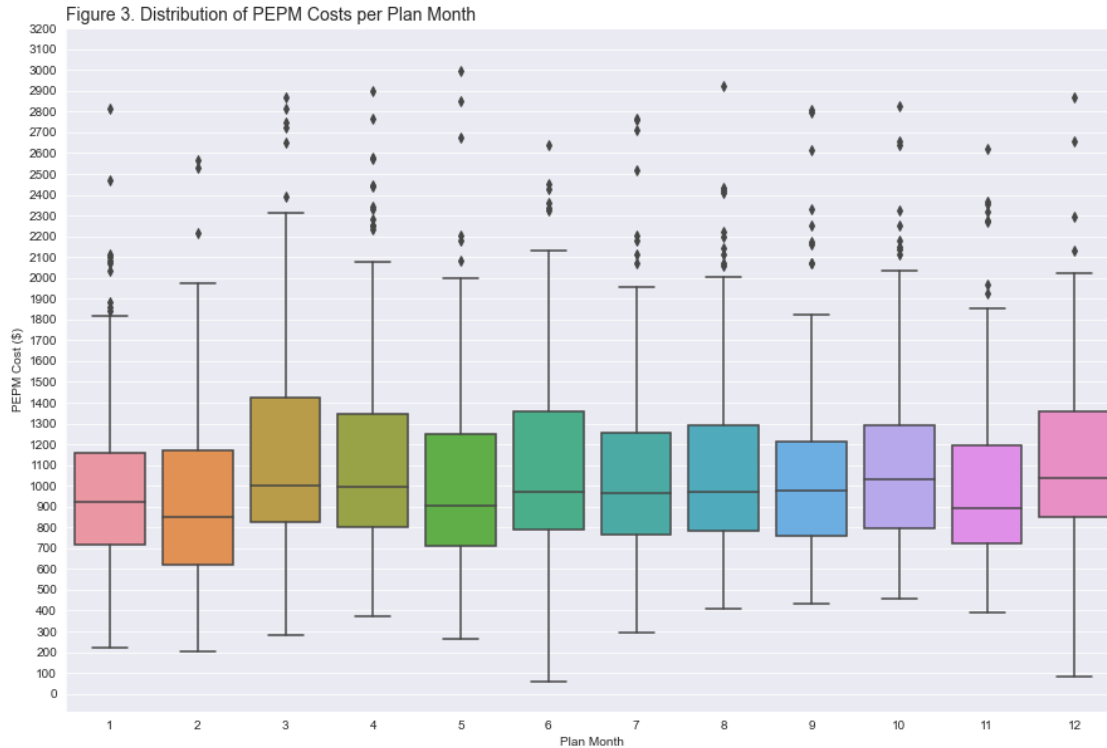



Figure 3 is particularly encouraging. In previous iterations of this project, a model was created with data just from Huntington Memorial Hospital. While the model worked, because of the extremely little data available, tempering volatility was an incredibly frustrating process. Here, while that same data is incorporated in the `claims` dataset, the sheer volume of data prevents outliers from hampering the data's utility. As a result, the mean of each `Plan Month` PEPM Cost distribution is relatively stable with similarly stable IQRs (middle 50%). Most of the outliers occur on the higher end which is to be expected.

```
[7]: # Figure 4. Overlaid Distributions of PEPM Cost per Plan Year Cohort
claims_copy = claims.copy()
claims_copy['Cohort'] = [2 if year >= 5 else 1 for year in claims_copy['Plan_
    ↳Year'].values]
plt.figure(figsize=(10, 5))
sns.histplot(data=claims_copy.query('PEPM > 0 and PEPM < 3000'), x='PEPM',
    ↳hue='Cohort', palette=['cornflowerblue', 'mediumseagreen'])
plt.xlabel('PEPM Cost ($)')
plt.title('Figure 4. Overlaid Distributions of PEPM Costs per Plan Year_
    ↳Cohort', loc='left', fontdict={'fontsize': 14})
plt.show()
cohort1 = claims_copy.query('Cohort == 1')['PEPM'].values
cohort2 = claims_copy.query('Cohort == 2')['PEPM'].values
score = stats.ttest_ind(cohort2, cohort1)
print(f'Statistic: {score[0]}')
```

```

print(f'p-value: {score[1]}')
print(f'Cohort 1 Proportion: {np.round(claims_copy.query("Cohort == 1").
    ↳shape[0]/claims.shape[0], 3)}')
print(f'Cohort 2 Proportion: {np.round(claims_copy.query("Cohort == 2").
    ↳shape[0]/claims.shape[0], 3)}')

```



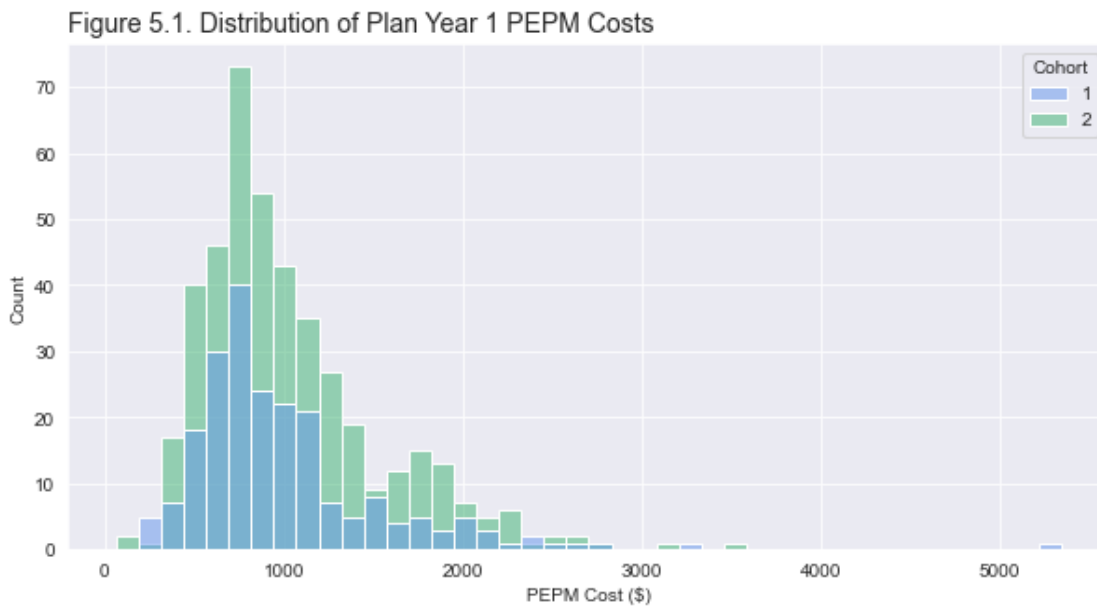
Statistic: 4.298020532918551
 p-value: 1.8288907954969783e-05
 Cohort 1 Proportion: 0.533
 Cohort 2 Proportion: 0.467

Like before, Figure 4 is also encouraging. The distribution of Cohort 1's (rows/data from Plan Years 1, 2, and 3) PEPM Costs is visually to the left (less) than that of Cohort 2 (rows/data from Plan Years 4 and later). This is backed up by the T-Test run on the two cohorts. From a high level overview, the Statistic value indicates how much greater the mean the distribution of Cohort 2 PEPM Costs is than that of Cohort 1 while the p-value indicates the level of certainty that this was not the result of pure chance (with a value closer to 0 indicative of a statistically significant difference). Given the rising costs of medical care, this is to be expected. As indicated by the `print` statements following the plot, it is also evident that the two cohorts contain approximately the same amount of data. The next thing to look at is whether cohorts can be created to approximate the given distributions of Plan Years 1, 2, 3, and 4 (or some other sequential order of four consecutive years).

The idea here is to find a cohort (defined as a subset of the `claims` dataset filtered by `Plan Year`) whose distribution is approximately the same as the distribution for each Plan Year to increase the sample size to randomly sample from in the Model Training/Modelling Stage.

```
[8]: # Figure 5.1. Distribution of Plan Year 1 PEPM Costs
claims_py1 = claims.copy()
claims_py1['Cohort'] = [1 if py == 1 else 2 if py <= 3 else 10 for py in
    ↪claims_py1['Plan Year'].values]
plt.figure(figsize=(10, 5))
sns.histplot(data=claims_py1.query('PEPM > 0 and Cohort < 3'), x='PEPM',
    ↪hue='Cohort', palette=['cornflowerblue', 'mediumseagreen'])
plt.xlabel('PEPM Cost ($)')
plt.title('Figure 5.1. Distribution of Plan Year 1 PEPM Costs', loc='left',
    ↪fontdict={'fontsize': 14})
plt.show()

cohort1 = claims_py1.query('Cohort == 1')['PEPM'].values
cohort2 = claims_py1.query('Cohort == 2')['PEPM'].values
score = stats.ttest_ind(cohort2, cohort1)
print(f'Statistic: {score[0]}')
print(f'p-value: {score[1]}')
```



Statistic: 0.6872124860475861
p-value: 0.4921954680866031

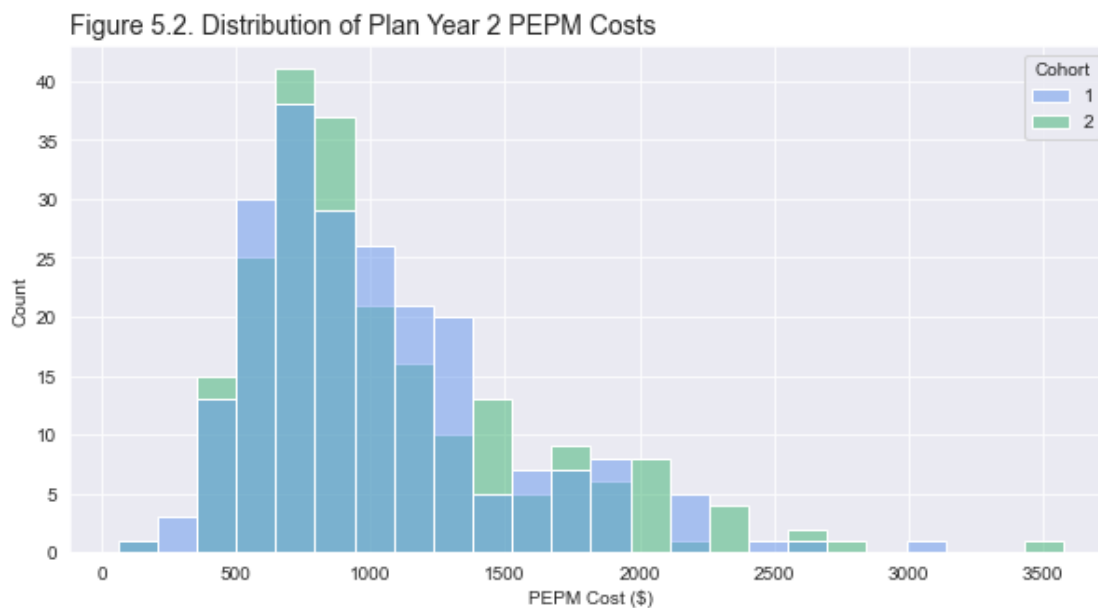
```
[9]: # Plan Year 2 Distribution of PEPM Costs
claims_py2 = claims.copy().query('`Plan Year` > 1')
claims_py2['Cohort'] = [1 if py == 2 else 2 if py <= 3 else 10 for py in
    ↪claims_py2['Plan Year'].values]
plt.figure(figsize=(10, 5))
```

```

sns.histplot(data=claims_py2.query('PEPM > 0 and Cohort < 3'), x='PEPM',
             hue='Cohort', palette=['cornflowerblue', 'mediumseagreen'])
plt.xlabel('PEPM Cost ($)')
plt.title('Figure 5.2. Distribution of Plan Year 2 PEPM Costs', loc='left',
         fontdict={'fontsize': 14})
plt.show()

cohort1 = claims_py2.query('Cohort == 1')['PEPM'].values
cohort2 = claims_py2.query('Cohort == 2')['PEPM'].values
score = stats.ttest_ind(cohort2, cohort1)
print(f'Statistic: {score[0]}')
print(f'p-value: {score[1]}')

```



Statistic: 0.9793483353514498
p-value: 0.32795856424709857

```

[10]: # Plan Year 3 Distribution of PEPM Costs
claims_py3 = claims.copy().query('`Plan Year` > 2')
claims_py3['Cohort'] = [1 if py == 3 else 2 if py <= 6 else 10 for py in
                        claims_py3['Plan Year'].values]
plt.figure(figsize=(10, 5))
sns.histplot(data=claims_py3.query('PEPM > 0 and Cohort < 3'), x='PEPM',
             hue='Cohort', palette=['cornflowerblue', 'mediumseagreen'])
plt.xlabel('PEPM Cost ($)')
plt.title('Figure 5.3. Distribution of Plan Year 3 PEPM Costs', loc='left',
         fontdict={'fontsize': 14})
plt.show()


```

```

cohort1 = claims_py3.query('Cohort == 1')['PEPM'].values
cohort2 = claims_py3.query('Cohort == 2')['PEPM'].values
score = stats.ttest_ind(cohort2, cohort1)
print(f'Statistic: {score[0]}')
print(f'p-value: {score[1]}')

```

Figure 5.3. Distribution of Plan Year 3 PEPM Costs




Statistic: 1.3089034290073263
p-value: 0.19098935580093418

```

[11]: # Plan Year 4 Distribution of PEPM Costs
claims_py4 = claims.copy().query('`Plan Year` > 3')
claims_py4['Cohort'] = [1 if py == 4 else 2 if py <= 8 else 10 for py in
    ↪claims_py4['Plan Year'].values]
plt.figure(figsize=(10, 5))
sns.histplot(data=claims_py4.query('PEPM > 0 and Cohort < 3'), x='PEPM',
    ↪hue='Cohort', palette=['cornflowerblue', 'mediumseagreen'])
plt.xlabel('PEPM Cost ($)')
plt.title('Figure 5.4. Distribution of Plan Year 4 PEPM Costs', loc='left',
    ↪fontdict={'fontsize': 14})
plt.show()

cohort1 = claims_py4.query('Cohort == 1')['PEPM'].values
cohort2 = claims_py4.query('Cohort == 2')['PEPM'].values
score = stats.ttest_ind(cohort2, cohort1)
print(f'Statistic: {score[0]}')
print(f'p-value: {score[1]}')

```



Statistic: 0.0907977145829931

p-value: 0.9276780306828769

Figures 5.1, 5.2, 5.3, and 5.4 are all analagous to Figure 4 above. While the p-value for Figure 5.3 is particularly lower than its peers, it does not pose any problems. More specifically, the p-value indicates (as a probability) the chance that the observed difference in means between the two PEPM Cost distributions is incorrectly deemed statistically significant. While ~19% may sound worryingly high, the gold standard in statistics is a p-value of 0.05 (or a 5% chance). If the p-value falls below the threshold, the phenomenon is generally deemed statistically significant. Since the lowest p-value is still nearly four times more than the traditional p-value cutoff, it is safe to say that the distributions generated are good approximations of the distributions of Plan Years 1, 2, 3, and 4. Next, these approximate distributions will be used to create a larger dataset with which to train the model.

1.3 Modelling

1.3.1 Data Formation

```
[12]: # Python user-defined functions in order of appearance
def resample_plan_year_cohort(cohort_df, n):
    """
    Generates resampled_plan_year_cohort DataFrame with n rows per month (i.e. 12n rows total)
    sampled with replacement from DataFrame cohort_df
    -----
    Inputs:
```

```

    cohort_df (DataFrame) - DataFrame with approximate distribution and
    ↳ filtered to exclude
                                negative and 0 values for PEPM Cost

    n (int) - Specifies the number of resamples
    -----
    Output:

    resampled_plan_year_cohort (DataFrame) - DataFrame generated from sampling
    ↳ with replacement
                                                with n rows per month (for a total
    ↳ of 12n rows)
'''
resampled_plan_year_cohort = pd.DataFrame(columns=cohort_df.columns)
for month in np.arange(1, 13):
    month_df = cohort_df.query('`Plan Month` == @month')
    resampled = month_df.sample(n=n, replace=True, axis=0)
    resampled_plan_year_cohort = resampled_plan_year_cohort.
    ↳ append(resampled)
return resampled_plan_year_cohort

```

The process followed will be a sort of pseudo-bootstrapping process to help create a larger dataset (expanded_claims) to train the model. Traditionally, bootstrapping is used as a tool to help calculate summary statistics or conduct hypothesis tests of small yet still representative samples. The small sample is sampled with replacement numerous times and a statistic calculated. The calculated statistics are then plotted in a histogram and build a generally normal distribution (of the summary statistic). In this case, sampling with replacement numerous times will be used to craft a larger dataset from the approximate distributions found in the previous section. Since the approximate distributions are roughly normal and representative of the actual distribution of PEPM Costs for Plan Years 1, 2, 3, and 4, the process of sampling with replacement will aid in filling in the gaps left by the relatively small sample size while not compromising the original claims dataset.

```

[13]: # Plan Year 1 Cohort
plan_year1_cohort = claims.query('PEPM > 0 and `Plan Year` <= 3')
cohort1 = resample_plan_year_cohort(plan_year1_cohort, 1000)

# Plan Year 2 Cohort
plan_year2_cohort = claims.query('PEPM > 0 and `Plan Year` > 1 and `Plan Year`
    ↳ <= 3')
cohort2 = resample_plan_year_cohort(plan_year2_cohort, 1000)

# Plan Year 3 Cohort
plan_year3_cohort = claims.query('PEPM > 0 and `Plan Year` > 2 and `Plan Year`
    ↳ <= 6')
cohort3 = resample_plan_year_cohort(plan_year3_cohort, 1000)

```



```

# Plan Year 4 Cohort
plan_year4_cohort = claims.query('PEPM > 0 and `Plan Year` > 3 and `Plan Year`_
↳ <= 8')
cohort4 = resample_plan_year_cohort(plan_year4_cohort, 1000)

# Combining the four resampled cohorts
cohorts = [cohort1, cohort2, cohort3, cohort4]
expanded_claims = pd.DataFrame(columns=claims.columns)
for cohort in cohorts:
    expanded_claims = expanded_claims.append(cohort)
expanded_claims = expanded_claims.drop(['Month', 'Year', 'Plan Year'], axis=1)
plan_years = np.array([], dtype=int)
ones = np.ones(1000, dtype=int)
twos = ones + 1
threes = ones + 2
fours = ones + 3
for i in np.arange(0, 12):
    plan_month_group = np.array([], dtype=int)
    plan_month_group = np.append(plan_month_group, ones)
    plan_month_group = np.append(plan_month_group, twos)
    plan_month_group = np.append(plan_month_group, threes)
    plan_month_group = np.append(plan_month_group, fours)
    plan_years = np.append(plan_years, plan_month_group)
expanded_claims['Plan Year'] = plan_years
expanded_claims = expanded_claims[['Plan Month',
    'Plan Year',
    'Total Medical Claims',
    'Total Rx Claims',
    'Total Paid Claims',
    'Adjusted Total Expenses',
    'EE Count',
    'Member Count',
    'Client',
    'PEPM',
    'PMPM']].reset_index().drop('index', axis=1)
expanded_claims

```

```

[13]:
    Plan Month  Plan Year  Total Medical Claims  Total Rx Claims  \
0           1         1           4767.00        195529.47
1           1         1        1324039.89        328595.73
2           1         1        395033.06        138763.63
3           1         1        1263663.73        451547.90
4           1         1        183404.43         85855.07
...         ...         ...         ...         ...
47995        12         4        1777329.15        484099.49
47996        12         4         941860.47         73348.48
47997        12         4       11646247.11       2220784.24

```

47998	12	4	2297491.03	894204.68
47999	12	4	315682.44	42450.61

	Total Paid Claims	Adjusted Total Expenses	EE Count	Member Count	\
0	200296.47	3.365825e+05	1093	2710	
1	1652635.62	1.842506e+06	2051	4973	
2	533796.69	6.264064e+05	1108	2598	
3	1715211.63	1.931644e+06	3020	6623	
4	269259.50	3.253038e+05	976	2057	
...	
47995	2261428.64	2.547698e+06	2706	6728	
47996	1015208.95	1.093209e+06	381	836	
47997	13867031.35	1.443582e+07	16422	32659	
47998	3191695.71	3.672852e+06	3048	5731	
47999	358133.05	4.234171e+05	611	1433	

	Client	PEPM	PMPM
0	Antelope Valley	307.943705	124.200173
1	Enloe	898.345256	370.501934
2	Marshall	565.348718	241.111001
3	Huntington Memorial Hospital	639.617053	291.656878
4	CHA Hollywood Pres	333.303064	158.144769
...
47995	Enloe	941.499531	378.670887
47996	Northern Inyo	2869.313955	1307.665810
47997	Prime Healthcare	879.053854	442.016669
47998	CHLA	1205.003773	640.874455
47999	Dameron	692.990387	295.476013

[48000 rows x 11 columns]

Now that `expanded_claims` has been constructed, some verification is necessary to ensure that the assumptions made in previous sections are valid.

```
[14]: plt.figure(figsize=(20, 20))

# Figure 6.1. Distribution of Expanded PEPM Costs per Plan Month
plt.subplot(2, 2, 1)
sns.boxplot(data=expanded_claims.query('PEPM < 3000'), x='Plan Month', y='PEPM')
plt.yticks(np.arange(0, 3300, 100))
plt.ylabel('PEPM Cost ($)')
plt.title('Figure 6.1. Distribution of Expanded PEPM Costs per Plan Month', loc='left', fontdict={'fontsize': 14})

# Figure 6.2. Distribution of Original PEPM Costs per Plan Month
plt.subplot(2, 2, 2)
```

```



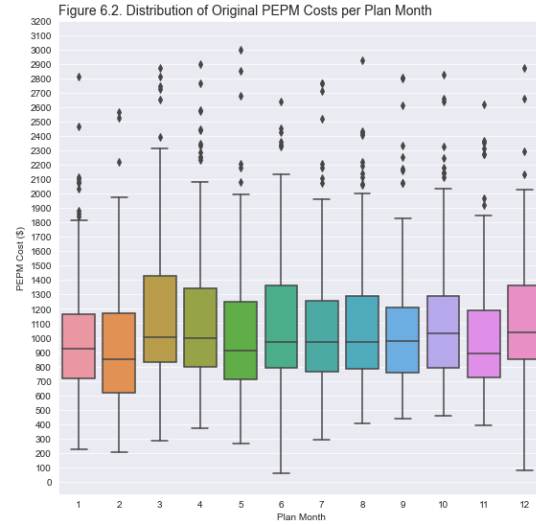
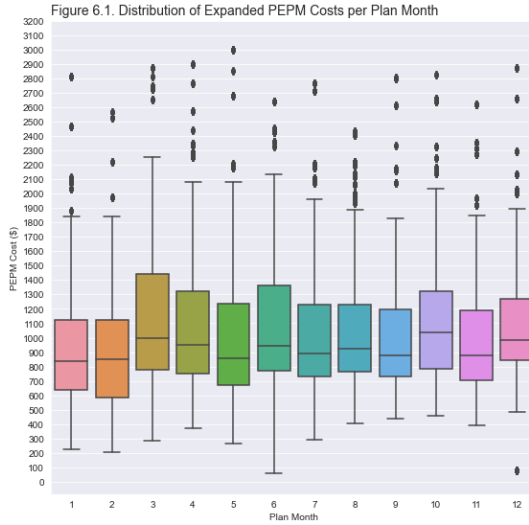
sns.boxplot(data=claims.query('PEPM > 0 and PEPM < 3000'), x='Plan Month',
            y='PEPM')
plt.yticks(np.arange(0, 3300, 100))
plt.ylabel('PEPM Cost ($)')
plt.title('Figure 6.2. Distribution of Original PEPM Costs per Plan Month',
            loc='left', fontdict={'fontsize': 14})

# Figure 6.3. Distribution of Expanded PMPM Costs per Plan Month
plt.subplot(2, 2, 3)
sns.boxplot(data=expanded_claims.query('PMPM < 1500'), x='Plan Month', y='PMPM')
plt.yticks(np.arange(0, 1700, 100))
plt.ylabel('PMPM Cost ($)')
plt.title('Figure 6.3. Distribution of Expanded PMPM Costs per Plan Month',
            loc='left', fontdict={'fontsize': 14})

# Figure 6.4. Distribution of Original PMPM Costs per Plan Month
plt.subplot(2, 2, 4)
sns.boxplot(data=claims.query('PMPM > 0 and PMPM < 1500'), x='Plan Month',
            y='PMPM')
plt.yticks(np.arange(0, 1700, 100))
plt.ylabel('PMPM Cost ($)')
plt.title('Figure 6.4. Distribution of Original PMPM Costs per Plan Month',
            loc='left', fontdict={'fontsize': 14})

plt.show()

```



These distributions look great. Comparing Figures 6.1 and 6.2 and Figures 6.3 and 6.4, the distributions look remarkably similar. This is precisely what was desired.

1.3.2 Transforming the Data

```
[15]: # Python user-defined functions in order of appearance
def calculate_year_avg_pepm(expanded_df, plan_year, n):
    """
    Extracts the year average PEPM cost from expanded DataFrame expanded_df
    for the specified plan year plan_year
    -----
    Input:

    resampled_df (DataFrame) - DataFrame with bootstrapped data
```

```

plan_year (int) - Int value specifying the renewal year to extract

n (int) - Number of resamples
-----
Output:

pepm_plan_year (NumPy Array) - NumPy Array of length n (where n is the
↳number of resamples)

                                with the yearly mean PEPM values
'''
plan_year = expanded_df.query('`Plan Year` == @plan_year')
pepm_arr_plan_year = plan_year['PEPM'].values
ee_arr_plan_year = plan_year['EE Count'].values
pepm_arr_plan_year_split = np.split(pepm_arr_plan_year, 12)
ee_arr_plan_year_split = np.split(ee_arr_plan_year, 12)
ee_sums = np.zeros(n)
for i in range(len(ee_arr_plan_year_split)):
    ee_sums = ee_sums + ee_arr_plan_year_split[i]
ee_props = [ee_arr_plan_year_split[i] / ee_sums for i in
↳range(len(ee_arr_plan_year_split))]
weighted_totals = [np.multiply(ee_props[i], pepm_arr_plan_year_split[i])
↳for i in range(len(ee_props))]
pepm_plan_year = np.zeros(n)
for i in range(len(weighted_totals)):
    pepm_plan_year = pepm_plan_year + weighted_totals[i]
return pepm_plan_year

def create_design_matrix(expanded_df, plan_year1, plan_year2, plan_year3, n,
↳for_training=True, PMPM=False):
    '''
    Creates a design matrix given the expanded DataFrame expanded_df,
    ↳plan year 1 DataFrame plan_year1,
    plan year 2 DataFrame plan_year2, and plan year 3 DataFrame plan_year3. If
    ↳specified, can also add PMPM
    features to the design matrix
    -----
    Inputs:

    expanded_df (DataFrame) - Expanded DataFrame used to calculate following
    ↳year average PEPM Cost

    plan_year1 (DataFrame) - Filtered expanded DataFrame with data from plan
    ↳year 1

```

```

    plan_year2 (DataFrame) - Filtered expanded DataFrame with data from plan_
↳year 2

    plan_year3 (DataFrame) - Filtered expanded DataFrame with data from plan_
↳year 3

    n (int) - Rows of data per month

    for_training (boolean) - Default value True; if False, creates design_
↳matrix without

                                Following Year Average PEPM column

    PMPM (boolean) - Default value False; if True includes PMPM value in design_
↳matrix matrix

    -----
    Output:

    matrix (DataFrame) - Design matrix as a DataFrame
    '''
    if for_training:
        actual_values = calculate_year_avg_pepm(expanded_df, 4, n)
        year1 = np.split(plan_year1['PEPM'].values, 12)
        year2 = np.split(plan_year2['PEPM'].values, 12)
        year3 = np.split(plan_year3['PEPM'].values, 12)
        if PMPM:
            year1_pmpm = np.split(plan_year1['PMPM'].values, 12)
            year2_pmpm = np.split(plan_year2['PMPM'].values, 12)
            year3_pmpm = np.split(plan_year3['PMPM'].values, 12)
            design_matrix = pd.DataFrame(data={'Month 1 PEPM': year1[0], 'Month 2_
↳PEPM': year1[1],
                                'Month 3 PEPM': year1[2], 'Month 4_
↳PEPM': year1[3],
                                'Month 5 PEPM': year1[4], 'Month 6_
↳PEPM': year1[5],
                                'Month 7 PEPM': year1[6], 'Month 8_
↳PEPM': year1[7],
                                'Month 9 PEPM': year1[8], 'Month 10_
↳PEPM': year1[9],
                                'Month 11 PEPM': year1[10], 'Month_
↳12 PEPM': year1[11],
                                'Month 13 PEPM': year2[0], 'Month 14_
↳PEPM': year2[1],
                                'Month 15 PEPM': year2[2], 'Month 16_
↳PEPM': year2[3],
                                'Month 17 PEPM': year2[4], 'Month 18_
↳PEPM': year2[5],

```

```

↪PEPM': year2[7],
↪PEPM': year2[9],
↪24 PEPM': year2[11],
↪PEPM': year3[1],
↪PEPM': year3[3],
↪PEPM': year3[5],
↪PEPM': year3[7],
↪PEPM': year3[9],
↪36 PEPM': year3[11],
↪'Month 2 PMPM': year1_pmpm[1],
↪'Month 4 PMPM': year1_pmpm[3],
↪'Month 6 PMPM': year1_pmpm[5],
↪'Month 8 PMPM': year1_pmpm[7],
↪'Month 10 PMPM': year1_pmpm[9],
↪'Month 12 PMPM': year1_pmpm[11],
↪'Month 14 PMPM': year2_pmpm[1],
↪'Month 16 PMPM': year2_pmpm[3],
↪'Month 18 PMPM': year2_pmpm[5],
↪'Month 20 PMPM': year2_pmpm[7],
↪'Month 22 PMPM': year2_pmpm[9],
↪'Month 24 PMPM': year2_pmpm[11],
↪'Month 26 PMPM': year3_pmpm[1],

'Month 19 PEPM': year2[6], 'Month 20_
'Month 21 PEPM': year2[8], 'Month 22_
'Month 23 PEPM': year2[10], 'Month_
'Month 25 PEPM': year3[0], 'Month 26_
'Month 27 PEPM': year3[2], 'Month 28_
'Month 29 PEPM': year3[4], 'Month 30_
'Month 31 PEPM': year3[6], 'Month 32_
'Month 33 PEPM': year3[8], 'Month 34_
'Month 35 PEPM': year3[10], 'Month_
'Month 1 PMPM': year1_pmpm[0],_
'Month 3 PMPM': year1_pmpm[2],_
'Month 5 PMPM': year1_pmpm[4],_
'Month 7 PMPM': year1_pmpm[6],_
'Month 9 PMPM': year1_pmpm[8],_
'Month 11 PMPM': year1_pmpm[10],_
'Month 13 PMPM': year2_pmpm[0],_
'Month 15 PMPM': year2_pmpm[2],_
'Month 17 PMPM': year2_pmpm[4],_
'Month 19 PMPM': year2_pmpm[6],_
'Month 21 PMPM': year2_pmpm[8],_
'Month 23 PMPM': year2_pmpm[10],_
'Month 25 PMPM': year3_pmpm[0],_

```

```

↪ 'Month 28 PMPM': year3_pmpm[3],
↪ 'Month 30 PMPM': year3_pmpm[5],
↪ 'Month 32 PMPM': year3_pmpm[7],
↪ 'Month 34 PMPM': year3_pmpm[9],
↪ 'Month 36 PMPM': year3_pmpm[11],
↪ actual_values.astype(float))
    else:
        design_matrix = pd.DataFrame(data={'Month 1': year1[0], 'Month 2': ↪
↪ year1[1],
↪ year1[3],
↪ year1[5],
↪ year1[7],
↪ year1[9],
↪ year1[11],
↪ year2[1],
↪ year2[3],
↪ year2[5],
↪ year2[7],
↪ year2[9],
↪ year2[11],
↪ year3[1],
↪ year3[3],
↪ year3[5],
↪ year3[7],
↪ 'Month 27 PMPM': year3_pmpm[2], ↪
↪ 'Month 29 PMPM': year3_pmpm[4], ↪
↪ 'Month 31 PMPM': year3_pmpm[6], ↪
↪ 'Month 33 PMPM': year3_pmpm[8], ↪
↪ 'Month 35 PMPM': year3_pmpm[10], ↪
↪ 'Following Year Average PEPM': ↪
↪ 'Month 3': year1[2], 'Month 4': ↪
↪ 'Month 5': year1[4], 'Month 6': ↪
↪ 'Month 7': year1[6], 'Month 8': ↪
↪ 'Month 9': year1[8], 'Month 10': ↪
↪ 'Month 11': year1[10], 'Month 12': ↪
↪ 'Month 13': year2[0], 'Month 14': ↪
↪ 'Month 15': year2[2], 'Month 16': ↪
↪ 'Month 17': year2[4], 'Month 18': ↪
↪ 'Month 19': year2[6], 'Month 20': ↪
↪ 'Month 21': year2[8], 'Month 22': ↪
↪ 'Month 23': year2[10], 'Month 24': ↪
↪ 'Month 25': year3[0], 'Month 26': ↪
↪ 'Month 27': year3[2], 'Month 28': ↪
↪ 'Month 29': year3[4], 'Month 30': ↪
↪ 'Month 31': year3[6], 'Month 32': ↪

```



```

        'Month 33': year3[8], 'Month 34':
    ↪year3[9],
        'Month 35': year3[10], 'Month 36':
    ↪year3[11],
        'Following Year Average PEPM':
    ↪actual_values.astype(float)})
    return design_matrix

```

Now that `expanded_claims` has been both created and validated, it is time to transform the data to the desired specs. Remember, the model is to take monthly PEPM/PMPM values to predict the following year's medical plan cost.

```

[16]: plan_year1 = expanded_claims.query('`Plan Year` == 1')
      plan_year2 = expanded_claims.query('`Plan Year` == 2')
      plan_year3 = expanded_claims.query('`Plan Year` == 3')
      design_matrix = create_design_matrix(expanded_claims, plan_year1, plan_year2,
    ↪plan_year3, 1000)
      design_matrix

```

```

[16]:
      Month 1      Month 2      Month 3      Month 4      Month 5  \
0      307.943705  1447.483855  620.953067  1104.748120  551.749646
1      898.345256  1447.483855  986.237859  653.828012  498.006966
2      565.348718  587.120072  1345.291253  1221.549597  722.779376
3      639.617053  738.483945  476.514541  1221.549597  551.749646
4      333.303064  639.062188  620.953067  591.971806  1344.957322
..      ...
995    593.838481  806.249597  1782.853885  596.386460  806.249597
996    435.059322  551.749646  986.237859  653.828012  443.697151
997    1123.820680  915.769056  861.483801  1381.770637  290.505933
998    1839.863038  778.365457  729.662604  773.966918  498.006966
999    1104.748120  778.365457  513.894667  599.668661  770.744479

      Month 6      Month 7      Month 8      Month 9      Month 10  ...  \
0    1270.926984  435.059322  1344.957322  1809.908190  836.400013  ...
1     484.485433  816.629079  1632.821306  1183.032702  537.818878  ...
2     902.911408  961.221509  332.492857  1061.825596  1459.345536  ...
3    1379.196353  1264.864778  820.982096  988.416138  882.876957  ...
4     654.236182  1150.542565  2998.296325  691.412192  816.629079  ...
..      ...
995    999.417022  925.451302  986.259779  786.898101  811.523028  ...
996    513.894667  812.119639  607.439074  794.672980  1346.174631  ...
997    513.894667  596.386460  986.259779  753.156211  878.245448  ...
998    513.894667  592.362742  1131.264403  833.288530  1150.542565  ...
999    805.285554  1817.434461  906.997174  1114.778080  1074.631429  ...

      Month 28      Month 29      Month 30      Month 31      Month 32  \
0    2225.839365  1476.196388  1291.047661  1059.306452  828.848517

```

1	533.108959	1958.491658	946.483291	810.266125	1795.830721
2	655.728431	799.438519	540.977703	1277.814544	765.371724
3	737.168251	1476.196388	2621.597415	1277.814544	2762.686675
4	776.618893	563.686714	1011.037134	750.510817	1175.415168
..
995	368.576364	550.461998	2621.597415	1084.994029	698.288445
996	2225.839365	678.306848	1491.320568	1925.291049	772.857322
997	533.108959	1848.786371	2274.115509	1188.918973	801.182338
998	927.726551	678.306848	901.663523	1072.231257	553.158903
999	462.950120	411.944739	660.007852	1863.809237	844.864065

	Month 33	Month 34	Month 35	Month 36	\
0	1015.711834	826.776137	1199.841702	695.419041	
1	691.819292	1594.189156	1454.776062	755.902907	
2	953.483619	1047.242055	929.851789	1203.401890	
3	725.742900	829.765179	929.851789	861.038773	
4	1093.058552	968.636367	1131.276837	731.276693	
..	
995	987.883676	665.924037	760.190761	731.276693	
996	1392.357859	956.088276	1467.074551	1036.574258	
997	1121.850965	1708.238068	649.132236	1040.211610	
998	1491.320568	826.776137	872.096832	826.823600	
999	1571.870538	800.603988	1680.665558	691.819292	

	Following Year Average PEPM
0	849.905109
1	1214.921388
2	1062.653172
3	1158.697167
4	1035.048479
..	...
995	902.315540
996	960.150204
997	946.205227
998	1035.201755
999	852.369119

[1000 rows x 37 columns]

Now that a design matrix has been created it can be used to train the model.

1.3.3 Training the Model

```
[17]: # Python user-defined functions in order of appearance
def create_test_matrix(data, first_py, PMPM=False):
    """
```

Given DataFrame data with three Plan Years worth of data, creates a test_
 ↪matrix to use for prediction

Input:

data (DataFrame) - DataFrame with three Plan Years worth of data

first_py (int) - Indicates which plan year is the first

*PMPM (Boolean) - Default value False; indicates whether PMPM should be_
 ↪included*

Output:

test_matrix (DataFrame) - Test matrix DataFrame for prediction
 '''

```

year1 = np.split(data.query('`Plan Year` == @first_py')['PEPM'].values, 12)
year2 = np.split(data.query('`Plan Year` == @first_py + 1')['PEPM'].values, ↪
↪12)
year3 = np.split(data.query('`Plan Year` == @first_py + 2')['PEPM'].values, ↪
↪12)
test_matrix = pd.DataFrame(data={'Month 1': year1[0], 'Month 2': year1[1],
                                'Month 3': year1[2], 'Month 4': year1[3],
                                'Month 5': year1[4], 'Month 6': year1[5],
                                'Month 7': year1[6], 'Month 8': year1[7],
                                'Month 9': year1[8], 'Month 10': year1[9],
                                'Month 11': year1[10], 'Month 12': ↪
↪year1[11],
                                'Month 13': year2[0], 'Month 14': year2[1],
                                'Month 15': year2[2], 'Month 16': year2[3],
                                'Month 17': year2[4], 'Month 18': year2[5],
                                'Month 19': year2[6], 'Month 20': year2[7],
                                'Month 21': year2[8], 'Month 22': year2[9],
                                'Month 23': year2[10], 'Month 24': ↪
↪year2[11],
                                'Month 25': year3[0], 'Month 26': year3[1],
                                'Month 27': year3[2], 'Month 28': year3[3],
                                'Month 29': year3[4], 'Month 30': year3[5],
                                'Month 31': year3[6], 'Month 32': year3[7],
                                'Month 33': year3[8], 'Month 34': year3[9],
                                'Month 35': year3[10], 'Month 36': ↪
↪year3[11]})
if PMPM:
    year1_pmpm = np.split(data.query('`Plan Year` == @first_py')['PMPM'].
↪values, 12)

```

```

year2_pmpm = np.split(data.query('`Plan Year` == @first_py +
↪1')['PMPM'].values, 12)
year3_pmpm = np.split(data.query('`Plan Year` == @first_py +
↪2')['PMPM'].values, 12)
test_matrix = pd.DataFrame(data={'Month 1 PEPM': year1[0], 'Month 2_
↪PEPM': year1[1],
                                'Month 3 PEPM': year1[2], 'Month 4_
↪PEPM': year1[3],
                                'Month 5 PEPM': year1[4], 'Month 6_
↪PEPM': year1[5],
                                'Month 7 PEPM': year1[6], 'Month 8_
↪PEPM': year1[7],
                                'Month 9 PEPM': year1[8], 'Month 10_
↪PEPM': year1[9],
                                'Month 11 PEPM': year1[10], 'Month 12_
↪PEPM': year1[11],
                                'Month 13 PEPM': year2[0], 'Month 14_
↪PEPM': year2[1],
                                'Month 15 PEPM': year2[2], 'Month 16_
↪PEPM': year2[3],
                                'Month 17 PEPM': year2[4], 'Month 18_
↪PEPM': year2[5],
                                'Month 19 PEPM': year2[6], 'Month 20_
↪PEPM': year2[7],
                                'Month 21 PEPM': year2[8], 'Month 22_
↪PEPM': year2[9],
                                'Month 23 PEPM': year2[10], 'Month 24_
↪PEPM': year2[11],
                                'Month 25 PEPM': year3[0], 'Month 26_
↪PEPM': year3[1],
                                'Month 27 PEPM': year3[2], 'Month 28_
↪PEPM': year3[3],
                                'Month 29 PEPM': year3[4], 'Month 30_
↪PEPM': year3[5],
                                'Month 31 PEPM': year3[6], 'Month 32_
↪PEPM': year3[7],
                                'Month 33 PEPM': year3[8], 'Month 34_
↪PEPM': year3[9],
                                'Month 35 PEPM': year3[10], 'Month 36_
↪PEPM': year3[11],
                                'Month 1 PMPM': year1_pmpm[0], 'Month_
↪2 PMPM': year1_pmpm[1],
                                'Month 3 PMPM': year1_pmpm[2], 'Month_
↪4 PMPM': year1_pmpm[3],

```

```

↪6 PMPM': year1_pmpm[5],
↪8 PMPM': year1_pmpm[7],
↪10 PMPM': year1_pmpm[9],
↪'Month 12 PMPM': year1_pmpm[11],
↪14 PMPM': year2_pmpm[1],
↪16 PMPM': year2_pmpm[3],
↪18 PMPM': year2_pmpm[5],
↪20 PMPM': year2_pmpm[7],
↪22 PMPM': year2_pmpm[9],
↪'Month 24 PMPM': year2_pmpm[11],
↪26 PMPM': year3_pmpm[1],
↪28 PMPM': year3_pmpm[3],
↪30 PMPM': year3_pmpm[5],
↪32 PMPM': year3_pmpm[7],
↪34 PMPM': year3_pmpm[9],
↪'Month 36 PMPM': year3_pmpm[11]})
    return test_matrix

def prediction(data, classification, under, over, verbose=False):
    """
    Given a singular row of 36 months of data (data), will return the predicted
    ↪yearly average PEPM
    cost for the following year following the procedure outlined in this report
    -----
    Inputs:

    data (DataFrame) - DataFrame (1, 36) with 36 months of scaled actual PEPM
    ↪values

```

```

    classification (sklearn.ensemble or sklearn.linear_model) - Sklearn
    ↳ classification model (either
                                                    Random Forest
    ↳ or a Logistic Regression
                                                    model) that
    ↳ predicts underestimates

    under (sklearn.linear_model or sklearn.neural_network) - Sklearn regression
    ↳ model (either Lasso Regression,
                                                    Ridge Regression,
    ↳ Linear Regression, or MLP Neural
                                                    Network) that
    ↳ predicts the average PEPM cost for the
                                                    following year for
    ↳ underestimated data

    over (sklearn.linear_model or sklearn.neural_network) - Sklearn regression
    ↳ model (either Lasso Regression,
                                                    Ridge Regression,
    ↳ Linear Regression, or MLP Neural
                                                    Network) that
    ↳ predicts the average PEPM cost for the
                                                    following year for
    ↳ overestimated data

    Y_scaler (sklearn.preprocessing._data.StandardScaler) - Fitted sklearn
    ↳ StandardScaler object to unscale prediction

    verbose (Boolean) - Default False; if True, prints 1 if underestimate and 0
    ↳ if not
    -----
    Output:

    prediction (float) - Predicted average PEPM cost for the following year
    '''
    is_underestimate = classification.predict(data)
    if verbose:
        estimate = ['Underestimate' if value == 1 else 'Overestimate' for value
    ↳ in is_underestimate][0]
        print(estimate)
    if is_underestimate:
        prediction = under.predict(data)
    else:
        prediction = over.predict(data)
    return prediction

```

Considering the relatively large number of features included in the design matrix, a few models

may be appropriate. In this section, two different approaches will be taken: a linear regression approach and a neural network approach.

The first approach is a Ridge Regression model (or an Ordinary Least Squares model with an L2 regularization term). At a high level, this is a linear model that allows for better generalizability. Ultimately, the key to creating a good model is to strike the optimal balance between model accuracy (how well the model can predict from its training data) and model generalizability (how well the model can predict from data it has never seen before). With a large number of features, the model in this project is particularly prone to overfitting to the training data (sacrificing generalizability in the process). Ridge Regression seeks to remedy this.

For a more sophisticated explanation for those more familiar with linear regression, think about the process behind linear regression. Under the hood, an OLS model is optimizing parameters that minimize some loss function (squared loss in this case). Adding L2 regularization, in effect, adds a penalty term to the loss function that penalizes the model proportional to the model's relative complexity in the hopes of preventing overfitting. Some hyperparameter, alpha, can also be tuned to allow for greater control of the strength of the regularization term&em; with larger values of alpha corresponding to a stronger regularization term.

The second approach will be a Multi-Layer Perceptron (MLP) Neural Network model. An MLP Neural Network is far more complex than a Ridge Regression model and functions by mimicking the natural learning process in the brain. Fundamentally, the model is comprised of various layers of nodes or perceptrons with each node in one layer connected to each of the nodes in the subsequent layer. Between layers, the signals fed into the model (e.g. the various features) undergoes a series of transformations determined by a series of hyperparameters that ultimately map to a singular value—in this specific case at least.

Disclaimer: I have not taken a formal course on Deep Learning and Neural Networks and thus am not able to discuss them on a more in-depth basis. This is, however, a high level understanding of what they are and how they differ from linear models.

Below, both models will be trained and their performance evaluated against each other.

```
[18]: # Split into train and test sets
X = design_matrix.drop('Following Year Average PEPM', axis=1)
Y = design_matrix[['Following Year Average PEPM']]
X_train, X_test, Y_train, Y_test = model_selection.train_test_split(X, Y,
    ↳test_size=0.2)

# Train and fit the model
model = linear_model.RidgeCV(alphas=np.arange(1, 10001), fit_intercept=False,
    ↳cv=None)
model.fit(X_train, Y_train.values.flatten())

# Model Validation
training_predictions = model.predict(X_train)
validation_predictions = model.predict(X_test)
model_training_mse = np.mean((training_predictions - Y_train.values.flatten())
    ↳** 2)
```

```

model_validation_mse = np.mean((validation_predictions - Y_test.values.
    ↪flatten()) ** 2)
print(f'Training MSE: {model_training_mse}')
print(f'Validation MSE: {model_validation_mse}')
print(min(training_predictions.flatten()))
print(max(training_predictions.flatten()))
print(f'Alpha: {model.alpha_}')

# Plotting the residuals
residuals = training_predictions - Y_train.values.flatten()
residuals_df = pd.DataFrame(data={'Predicted Value': training_predictions,
    ↪'Residual': residuals}, dtype=float)

plt.figure(figsize=(15, 5))
plt.subplot(1, 2, 1)
sns.scatterplot(data=residuals_df, x='Predicted Value', y='Residual', alpha=0.
    ↪25)
plt.title('Figure 7.1. Ridge Regression Model Residual Plot', loc='left',
    ↪fontdict={'fontsize': 14})
plt.show()
correlation = residuals_df['Predicted Value'].corr(residuals_df['Residual'])
print(f'Correlation: {correlation}')

# Plotting model parameters
model_parameters = model.coef_.flatten()
coefficients_df = pd.DataFrame(data={'Month': np.arange(1, 37), 'Coefficients':
    ↪model_parameters})

plt.figure(figsize=(15, 7.5))
sns.barplot(data=coefficients_df, x='Month', y='Coefficients')
plt.xlabel('Month')
plt.ylabel('Model Parameter')
plt.title('Figure 7.2. Ridge Regression Model Parameter By Month', loc='left',
    ↪fontdict={'fontsize': 14})
plt.show()

# Train and fit a MLP Neural Network
neural_net = neural_network.MLPRegressor(hidden_layer_sizes=(36,18,9),
    activation='relu',
    alpha=0.001,
    max_iter=1000,
    early_stopping=True)
neural_net.fit(X_train, Y_train.values.flatten())

# Neural Network Validation
neural_network_training_predictions = neural_net.predict(X_train)

```



```

neural_network_validation_predictions = neural_net.predict(X_test)
neural_network_training_mse = np.mean((neural_network_training_predictions -
    ↪ Y_train.values.flatten()) ** 2)
neural_network_validation_mse = np.mean((neural_network_validation_predictions -
    ↪ Y_test.values.flatten()) ** 2)
print(f'Training MSE: {neural_network_training_mse}')
print(f'Validation MSE: {neural_network_validation_mse}')
print(min(neural_network_training_predictions.flatten()))
print(max(neural_network_training_predictions.flatten()))

# Plotting the residuals
neural_network_residuals = neural_network_training_predictions - Y_train.values.
    ↪ flatten()
neural_network_residuals_df = pd.DataFrame(data={'Predicted Value':
    ↪ neural_network_training_predictions,
                                                'Residual':
    ↪ neural_network_residuals}, dtype=float)

plt.figure(figsize=(15, 5))
plt.subplot(1, 2, 1)
sns.scatterplot(data=neural_network_residuals_df, x='Predicted Value',
    ↪ y='Residual', alpha=0.25)
plt.title('Figure 7.3. MLP Neural Network Residual Plot', loc='left',
    ↪ fontdict={'fontsize': 14})
plt.show()
neural_network_correlation = neural_network_residuals_df['Predicted Value'] \
    .corr(neural_network_residuals_df['Residual'])
print(f'Correlation: {neural_network_correlation}')

```


Training MSE: 22229.717116708485
 Validation MSE: 21162.555018187468
 765.2750230512275
 1227.6825774250406
 Alpha: 10000

Figure 7.1. Ridge Regression Model Residual Plot




Correlation: 0.47505801588270896

Figure 7.2. Ridge Regression Model Parameter By Month



Training MSE: 24734.121436783345
Validation MSE: 25691.442069255645
674.2621944787977
1449.8383842432252

Figure 7.3. MLP Neural Network Residual Plot



Correlation: 0.5603363834998542

Figures 7.1 and 7.3 plot the residuals (or errors of the predicted values) against the predicted value itself. These plots are critical to interpreting the quality of a regression model.

Looking to both plots it is clear that both models have residuals that—while centered around zero—have a clear trend. With negative residuals, this indicates that both models have a tendency to underestimate values. To remedy this, a classification model can be utilized to identify datapoints that may be more prone to over/underestimation which may can be used to train a separate model to make better predictions for these data. Since the Ridge Regression model had better performance (lower MSE), this model will be used to build training data for a new classification model.

```
[19]: # Add Underestimate label to training data
estimation = X_train.copy()
estimation['Underestimate'] = [1 if residual < 0 else 0 for residual in ↵
    ↵residuals_df['Residual'].values]
estimation
```

[19]:

	Month 1	Month 2	Month 3	Month 4	Month 5	\
235	1747.065226	673.225215	850.588610	594.536481	1391.845173	
600	756.333005	262.300574	727.682069	639.617053	673.225215	
355	2103.206342	867.366382	1289.258013	2032.877127	1277.835815	
778	773.966918	1238.367581	1379.196353	538.555826	603.557773	
247	542.571954	1344.957322	476.514541	2467.401134	711.889838	
..	
849	538.555826	658.930823	620.953067	1104.748120	1510.190543	
574	773.165335	498.006966	629.462474	1118.556733	443.697151	
742	1882.033827	722.779376	1495.060530	653.828012	1154.850060	
609	1159.309037	1510.190543	476.514541	435.059322	443.697151	
288	565.348718	641.503610	1098.713400	812.119639	1298.647023	
	Month 6	Month 7	Month 8	Month 9	Month 10	...
235	999.417022	797.842914	1781.540192	761.892292	898.465102	...
600	861.483801	811.382128	1131.264403	761.892292	833.578781	...
355	986.237859	1729.247637	1391.845173	1480.810411	592.362742	...
778	1270.926984	773.966918	915.783403	1276.993733	2814.879339	...
247	2796.614031	882.876957	1382.918939	1111.880839	1175.745104	...
..
849	667.416984	757.461786	948.213252	1266.194484	592.946100	...
574	563.919551	811.382128	1077.907728	649.853196	1459.345536	...
742	784.180990	550.186267	747.605751	753.156211	1089.978008	...
609	999.417022	905.351070	640.126646	756.453981	1459.345536	...
288	1098.713400	728.053144	2179.617059	691.412192	1346.174631	...
	Month 28	Month 29	Month 30	Month 31	Month 32	\
235	1855.998373	1162.670475	391.061144	508.004663	1904.515471	
600	1072.231257	1162.670475	1093.058552	999.185395	828.848517	
355	1510.145745	811.730327	727.317420	1133.207724	811.730327	
778	607.674018	678.306848	841.659752	1133.207724	1199.841702	
247	761.120391	1476.196388	840.363715	810.266125	1795.830721	
..	
849	462.950120	591.305452	1291.047661	815.057577	1127.859191	
574	1796.991864	924.121196	901.663523	900.412461	2762.686675	
742	368.576364	550.461998	727.317420	855.872918	760.190761	
609	1438.646010	1737.730464	552.941686	861.997688	847.481353	
288	999.185395	891.917518	841.659752	1209.128083	1127.859191	
	Month 33	Month 34	Month 35	Month 36	Underestimate	
235	987.883676	1669.309787	998.583770	881.966942		0
600	712.839782	1209.128083	839.386056	615.532467		1
355	739.172328	826.776137	1231.275516	588.318148		0
778	691.819292	861.997688	828.848517	826.823600		1
247	691.819292	1277.814544	802.956574	1203.309866		1
..	
849	953.483619	884.598194	695.050192	831.983866		1

574	849.653799	884.598194	1009.920272	731.276693	1
742	2270.261938	1084.906065	839.386056	543.060827	0
609	657.674832	829.765179	984.601840	831.983866	1
288	885.453551	884.598194	1904.515471	1541.978883	0

[800 rows x 37 columns]

Using this new data **estimation**, a classification model can be trained to identify data prone to overestim. Two different classification models will be trained to determine the best approach: a Logistic Regression model and a Random Forest.

Although regression is in the name, Logistic Regression models are used primarily for binary classification. At a high level, Logistic Regression models predict the probability that a datapoint belongs to one label over the other. If the calculated probability is greater than some threshold (typically 0.5), the model predicts it belongs to label 1 and if not to label 0.

On the other hand, a Random Forest utilizes Decision Trees to classify data. Fundamentally, Decision Trees split data by some criteria (e.g. the value of a specific feature) until the training data is as pure as possible. A Random Forest takes a number of Decision Trees and has them “vote” to determine the label. Each Decision Tree is trained with a handle of the training data provided. The hope is to intentionally introduce randomness to increase the variance of individual Decision Trees. By combining many Decision Trees, a Random Forest (what is known as an ensemble model) hopes to account for that randomness.

Both will be trained and evaluated below.

```
[20]: # Split into train and test sets
X_estimation = estimation.drop('Underestimate', axis=1)
Y_estimation = estimation[['Underestimate']]
X_train2, X_test2, Y_train2, Y_test2 = model_selection.
    ↪train_test_split(X_estimation, Y_estimation, test_size=0.2)

# Train and fit the logistic regression model
classifier = linear_model.LogisticRegression(penalty='l2', solver='lbfgs',
    ↪fit_intercept=False, max_iter=10000)
classifier.fit(X_train2, Y_train2.values.flatten())

# Logistic Regression Validation
training_predictions2 = classifier.predict(X_train2)
validation_predictions2 = classifier.predict(X_test2)

print(f'Logistic Regression Training Accuracy: {classifier.score(X_train2,
    ↪Y_train2)}')
print(f'Logistic Regression Validation Accuracy: {classifier.score(X_test2,
    ↪Y_test2)}')

# Train and fit the random forest model
forest = ensemble.RandomForestClassifier(n_estimators=5000, max_depth=3,
    ↪min_samples_leaf=0.2, min_impurity_decrease=0.01, class_weight='balanced')
```

```

forest.fit(X_train2, Y_train2.values.flatten())

# Random Forest Validation
print(f'Random Forest Training Accuracy: {forest.score(X_train2, Y_train2)}')
print(f'Random Forest Validation Accuracy: {forest.score(X_test2, Y_test2)}')
print('\n')

# Confusion Matrices
print('Logistic Regression Confusion Matrix')
classifier_confusion = metrics.confusion_matrix(y_true=Y_test2.values.
    ↳flatten(), y_pred=validation_predictions2)
display(classifier_confusion)
classifier_precision = classifier_confusion[1][1]/(classifier_confusion[1][1] +
    ↳classifier_confusion[0][1])
print(f'Logistic Regression Precision: {classifier_precision}')
print('\n')
print('Random Forest Confusion Matrix')
forest_confusion = metrics.confusion_matrix(y_true=Y_test2.values.flatten(),
    ↳y_pred=forest.predict(X_test2))
display(forest_confusion)
forest_precision = forest_confusion[1][1]/(forest_confusion[1][1] +
    ↳forest_confusion[0][1])
print(f'Random Forest Precision: {forest_precision}')

```

```

Logistic Regression Training Accuracy: 0.553125
Logistic Regression Validation Accuracy: 0.40625
Random Forest Training Accuracy: 0.6828125
Random Forest Validation Accuracy: 0.55625

```

Logistic Regression Confusion Matrix

```

array([[45, 45],
       [50, 20]], dtype=int64)

```

Logistic Regression Precision: 0.3076923076923077

Random Forest Confusion Matrix

```

array([[58, 32],
       [39, 31]], dtype=int64)

```

Random Forest Precision: 0.49206349206349204

From the above, it is clear that the Random Forest model is the best suited for the task. Not only does it have the best validation accuracy, it also has higher precision—the proportion of predicted underestimates that are truly underestimates. This will ensure that the model will best split the data to isolate data prone to underestimation from that prone to overestimation. The predictions from this Random Forest model will be used to train a regression model for both the

underestimated data and the overestimated data. Like before, a Ridge Regression model and an MLP Neural Network model will be trained for both sets of data.

```
[21]: # Underestimate Modelling

underestimate_prediction_matrix = design_matrix.copy().drop('Following Year_
↳Average PEPM', axis=1)
underestimate_prediction_matrix['Underestimate'] = forest.
↳predict(underestimate_prediction_matrix)
underestimate_matrix = underestimate_prediction_matrix.copy()
underestimate_matrix['Following Year Average PEPM'] = design_matrix['Following_
↳Year Average PEPM']
underestimate_matrix = underestimate_matrix.query('Underestimate == 1').
↳reset_index().drop('index', axis=1) \
. drop('Underestimate', axis=1)

# Split into train and test sets
X_underestimate = underestimate_matrix.drop('Following Year Average PEPM',
↳axis=1)
Y_underestimate = underestimate_matrix[['Following Year Average PEPM']]
X_under_train, X_under_test, Y_under_train, Y_under_test = model_selection.
↳train_test_split(X_underestimate, Y_underestimate, test_size=0.2)

# Train and fit the Ridge Regression model
under_model = linear_model.RidgeCV(alphas=np.arange(1, 10001),
↳fit_intercept=False, cv=None)
under_model.fit(X_under_train, Y_under_train.values.flatten())

# Ridge Regression Model Validation
training_predictions_under = under_model.predict(X_under_train)
validation_predictions_under = under_model.predict(X_under_test)
under_model_training_mse = np.mean((training_predictions_under - Y_under_train.
↳values.flatten()) ** 2)
under_model_validation_mse = np.mean((validation_predictions_under -
↳Y_under_test.values.flatten()) ** 2)
print(f'Training Ridge Regression MSE: {under_model_training_mse}')
print(f'Validation Ridge Regression MSE: {under_model_validation_mse}')
print(min(training_predictions_under.flatten()))
print(max(training_predictions_under.flatten()))
print(f'Alpha: {under_model.alpha_}')

# Plotting the residuals
residuals_under = training_predictions_under - Y_under_train.values.flatten()
residuals_under_df = pd.DataFrame(data={'Predicted Value':
↳training_predictions_under, 'Residual': residuals_under}, dtype=float)

plt.figure(figsize=(15, 5))
```

```

plt.subplot(1, 2, 1)
sns.scatterplot(data=residuals_under_df, x='Predicted Value', y='Residual',
    ↪alpha=0.25)
plt.title('Figure 8.1. Underestimate Ridge Regression Model Residual Plot',
    ↪loc='left', fontdict={'fontsize': 14})
plt.show()
under_correlation = residuals_under_df['Predicted Value'].
    ↪corr(residuals_under_df['Residual'])
print(f'Correlation: {under_correlation}')

# Plotting Ridge Regression model parameters
under_model_parameters = under_model.coef_.flatten()
under_coefficients_df = pd.DataFrame(data={'Month': np.arange(1, 37),
    ↪'Coefficients': under_model_parameters})

plt.figure(figsize=(15, 7.5))
sns.barplot(data=under_coefficients_df, x='Month', y='Coefficients')
plt.xlabel('Month')
plt.ylabel('Model Parameter')
plt.title('Figure 8.2. Underestimate Ridge Regression Model Parameter By
    ↪Month', loc='left', fontdict={'fontsize': 14})
plt.show()

# Train MLP Neural Network for regression
under_mlp = neural_network.MLPRegressor(hidden_layer_sizes=(18,9),
    activation='relu',
    alpha=0.0000001,
    max_iter=1000,
    early_stopping=True)
under_mlp.fit(X_under_train, Y_under_train.values.flatten())

# MLP Neural Network model validation
nn_under_training_predictions = under_mlp.predict(X_under_train)
nn_under_validation_predictions = under_mlp.predict(X_under_test)
nn_under_model_training_mse = np.mean((nn_under_training_predictions -
    ↪Y_under_train.values.flatten()) ** 2)
nn_under_model_validation_mse = np.mean((nn_under_validation_predictions -
    ↪Y_under_test.values.flatten()) ** 2)
print(f'Training Neural Network MSE: {nn_under_model_training_mse}')
print(f'Validation Neural Network MSE: {nn_under_model_validation_mse}')
print(min(nn_under_training_predictions.flatten()))
print(max(nn_under_training_predictions.flatten()))

# Plotting the residuals
nn_residuals_under = nn_under_training_predictions - Y_under_train.values.
    ↪flatten()

```



```

nn_residuals_under_df = pd.DataFrame(data={'Predicted Value':  

    ↳nn_under_training_predictions, 'Residual': nn_residuals_under}, dtype=float)

plt.figure(figsize=(15, 5))
plt.subplot(1, 2, 1)
sns.scatterplot(data=nn_residuals_under_df, x='Predicted Value', y='Residual',  

    ↳alpha=0.25)
plt.title('Figure 8.3. Underestimate MLP Neural Network Residual Plot',  

    ↳loc='left', fontdict={'fontsize': 14})
plt.show()
nn_under_correlation = nn_residuals_under_df['Predicted Value'].  

    ↳corr(nn_residuals_under_df['Residual'])
print(f'Correlation: {nn_under_correlation}')

# Overestimation Modelling

overestimate_prediction_matrix = design_matrix.copy().drop('Following Year',  

    ↳Average PEPM', axis=1)
overestimate_prediction_matrix['Overestimate'] = forest.  

    ↳predict(overestimate_prediction_matrix)
overestimate_matrix = overestimate_prediction_matrix.copy()
overestimate_matrix['Following Year Average PEPM'] = design_matrix['Following  

    ↳Year Average PEPM']
overestimate_matrix = overestimate_matrix.query('Overestimate == 0').  

    ↳reset_index().drop('index', axis=1) \
        .drop('Overestimate', axis=1)

# Split into train and test sets
X_overestimate = overestimate_matrix.drop('Following Year Average PEPM', axis=1)
Y_overestimate = overestimate_matrix[['Following Year Average PEPM']]
X_over_train, X_over_test, Y_over_train, Y_over_test = model_selection.  

    ↳train_test_split(X_overestimate, Y_overestimate, test_size=0.2)

# Train and fit the Ridge Regression model
over_model = linear_model.RidgeCV(alphas=np.arange(1, 10001),  

    ↳fit_intercept=False, cv=None)
over_model.fit(X_over_train, Y_over_train.values.flatten())

# Ridge Regression Model Validation
training_predictions_over = over_model.predict(X_over_train)
validation_predictions_over = over_model.predict(X_over_test)
over_model_training_mse = np.mean((training_predictions_over - Y_over_train.  

    ↳values.flatten()) ** 2)
over_model_validation_mse = np.mean((validation_predictions_over - Y_over_test.  

    ↳values.flatten()) ** 2)
print(f'Training Ridge Regression MSE: {over_model_training_mse}')

```

```

print(f'Validation Ridge Regression MSE: {over_model_validation_mse}')
print(min(training_predictions_over.flatten()))
print(max(training_predictions_over.flatten()))
print(f'Alpha: {over_model.alpha_}')

# Plotting the residuals
residuals_over = training_predictions_over - Y_over_train.values.flatten()
residuals_over_df = pd.DataFrame(data={'Predicted Value': training_predictions_over, 'Residual': residuals_over}, dtype=float)

plt.figure(figsize=(15, 5))
plt.subplot(1, 2, 1)
sns.scatterplot(data=residuals_over_df, x='Predicted Value', y='Residual', alpha=0.25)
plt.title('Figure 8.4. Overestimate Ridge Regression Model Residual Plot', loc='left', fontdict={'fontsize': 14})
plt.show()
over_correlation = residuals_over_df['Predicted Value'].corr(residuals_over_df['Residual'])
print(f'Correlation: {over_correlation}')

# Plotting Ridge Regression model parameters
over_model_parameters = over_model.coef_.flatten()
over_coefficients_df = pd.DataFrame(data={'Month': np.arange(1, 37), 'Coefficients': over_model_parameters})

plt.figure(figsize=(15, 7.5))
sns.barplot(data=over_coefficients_df, x='Month', y='Coefficients')
plt.xlabel('Month')
plt.ylabel('Model Parameter')
plt.title('Figure 8.5. Overestimate Ridge Regression Model Parameter By Month', loc='left', fontdict={'fontsize': 14})
plt.show()

# Train MLP Neural Network for regression
over_mlp = neural_network.MLPRegressor(hidden_layer_sizes=(18,9),
                                       activation='relu',
                                       alpha=0.0000001,
                                       max_iter=1000,
                                       early_stopping=True)
over_mlp.fit(X_under_train, Y_under_train.values.flatten())

# MLP Neural Network model validation
nn_over_training_predictions = over_mlp.predict(X_over_train)
nn_over_validation_predictions = over_mlp.predict(X_over_test)

```

```

nn_over_model_training_mse = np.mean((nn_over_training_predictions -
    ↪Y_over_train.values.flatten()) ** 2)
nn_over_model_validation_mse = np.mean((nn_over_validation_predictions -
    ↪Y_over_test.values.flatten()) ** 2)
print(f'Overestimate Neural Network Training MSE: {nn_over_model_training_mse}')
print(f'Overestimate Neural Network Validation MSE:
    ↪{nn_over_model_validation_mse}')
print(min(nn_over_training_predictions.flatten()))
print(max(nn_over_training_predictions.flatten()))


# Plotting the residuals
nn_residuals_over = nn_over_training_predictions - Y_over_train.values.flatten()
nn_residuals_over_df = pd.DataFrame(data={'Predicted Value':
    ↪nn_over_training_predictions, 'Residual': nn_residuals_over}, dtype=float)

plt.figure(figsize=(15, 5))
plt.subplot(1, 2, 1)
sns.scatterplot(data=nn_residuals_over_df, x='Predicted Value', y='Residual',
    ↪alpha=0.25)
plt.title('Figure 8.6. Overestimate MLP Neural Network Residual Plot',
    ↪loc='left', fontdict={'fontsize': 14})
plt.show()
nn_over_correlation = nn_residuals_over_df['Predicted Value'].
    ↪corr(nn_residuals_over_df['Residual'])
print(f'Correlation: {nn_over_correlation}')

```


Training Ridge Regression MSE: 15203.70741077411
 Validation Ridge Regression MSE: 24988.549685380578
 747.1580598361725
 1250.375563824366
 Alpha: 10000

Figure 8.1. Underestimate Ridge Regression Model Residual Plot



Correlation: 0.356478023868552

Figure 8.2. Underestimate Ridge Regression Model Parameter By Month



Training Neural Network MSE: 15976.84557091129


Validation Neural Network MSE: 31078.0438098689
618.2167108696516
1208.054722503117

Figure 8.3. Underestimate MLP Neural Network Residual Plot




Correlation: 0.4260830258811177
Training Ridge Regression MSE: 20680.60725994785
Validation Ridge Regression MSE: 25819.70828254963
750.7965881994808
1267.4793906685998
Alpha: 10000

Figure 8.4. Overestimate Ridge Regression Model Residual Plot




Correlation: 0.39706927284210347

Figure 8.5. Overestimate Ridge Regression Model Parameter By Month



Overestimate Neural Network Training MSE: 101099.52121214487
Overestimate Neural Network Validation MSE: 102112.71709932554
179.77737318709384
1895.750293896821

Figure 8.6. Overestimate MLP Neural Network Residual Plot



Correlation: 0.9061885362473711

1.4 Validating the Model

Based on the validation MSE computed for each model from the previous section, the Ridge Regression models should be used. However, in practice, an argument can be made for the MLP Neural Network models as well. MSE is a measure of aggregate error. In some cases (shown below), the MLP Neural Network is far more accurate than the two Ridge Regression over and underestimation models and the original Ridge Regression model alone. It is the cases where the MLP Neural Network is dramatically off that drags its MSE lower than that of the over and underestimation Ridge Regression models. The models' performance on the historical `claims` data is shown below. See the Python comments for more details.

```
[22]: # Plan Year 4 Historical Predictions
client_list = claims.query('Client != "Epic" and Client !=_
↪ "Marshall")['Client'].unique()
```

```

first_three_list = [claims.query('Client == @client and (`Plan Year` == 1 or
↳ `Plan Year` == 2 or `Plan Year` == 3)') for client in client_list]
test_matrix_list = [create_test_matrix(client_data, 1) for client_data in
↳ first_three_list]
model_predictions = np.array([model.predict(test_matrix)[0] for test_matrix in
↳ test_matrix_list])
nn_prediction_list = np.array([prediction(test_matrix, forest, under_model,
↳ over_model)[0] for test_matrix in test_matrix_list])
prediction_list = np.array([prediction(test_matrix, forest, under_mlp,
↳ over_mlp)[0] for test_matrix in test_matrix_list])
actual_list = np.array([calculate_year_avg_pepm(claims.query('Client ==
↳ @client'), 4, 1)[0] for client in client_list])
py4_MSE = np.mean((prediction_list - actual_list) ** 2)
nn_py4_MSE = np.mean((nn_prediction_list - actual_list) ** 2)
original_py4_MSE = np.mean((model_predictions - actual_list) ** 2)
predictions = pd.DataFrame(data={'Predicted Values': prediction_list,
↳ 'Neural Network Predicted Values':
↳ nn_prediction_list,
↳ 'Original Model Predicted Values':
↳ model_predictions,
↳ 'Actual Values': actual_list,
↳ 'Client': client_list})
display(predictions)
print(f'Combined MSE: {py4_MSE}')
print(f'Neural Network MSE: {nn_py4_MSE}')
print(f'Original MSE: {original_py4_MSE}')

```

	Predicted Values	Neural Network Predicted Values \
0	948.181312	962.478837
1	700.034312	654.454889
2	844.343276	809.681250
3	560.915828	524.407371
4	1069.713506	968.599788
5	696.982554	669.066078
6	811.708483	797.303752
7	1943.057757	1699.458960
8	745.431231	832.273765
9	845.344197	796.644843
10	1512.900576	1627.887947
11	662.051496	616.277976
12	678.726388	618.085578
13	2035.463477	1609.794240
14	1184.108116	1179.747145
15	922.171479	878.908422

	Original Model Predicted Values	Actual Values \
0	1013.265594	1294.984380

1	609.788828	759.745909
2	776.753063	778.439771
3	501.032234	839.512584
4	1006.095069	1179.411289
5	637.789061	632.044792
6	831.265576	980.934774
7	1769.023737	2134.359135
8	846.833562	922.635962
9	759.191388	805.821466
10	1736.639784	2231.387449
11	583.727175	739.065477
12	561.555774	792.159896
13	1705.630825	1983.837721
14	1242.630714	1171.634353
15	847.709679	944.566036

	Client
0	Antelope Valley
1	Avanti/Pipeline
2	Beverly Hospital
3	CHA Hollywood Pres
4	CHLA
5	Dameron
6	Enloe
7	Fairchild
8	Henry Mayo
9	Huntington Memorial Hospital
10	Northern Inyo
11	Prime Healthcare
12	Prospect Medical Holdings
13	Salinas Valley
14	Tahoe Forest
15	Torrance Health

Combined MSE: 53661.82072883048
 Neural Network MSE: 65794.32060574666
 Original MSE: 51511.95672582195

```
[23]: # Plan Year 5 Historical Predictions
client_list2 = claims.query('Client != "Epic" and Client != "Marshall" and
    ↳ Client != "CHA Hollywood Pres" and Client != "Northern Inyo" and Client !=
    ↳ "Salinas Valley")['Client'].unique()
first_three_list = [claims.query('Client == @client and (`Plan Year` == 2 or
    ↳ `Plan Year` == 3 or `Plan Year` == 4)') for client in client_list2]
test_matrix_list = [create_test_matrix(client_data, 2) for client_data in
    ↳ first_three_list]
model_predictions = np.array([model.predict(test_matrix)[0] for test_matrix in
    ↳ test_matrix_list])
```

```

prediction_list = np.array([prediction(test_matrix, forest, under_model,
↳over_model)[0] for test_matrix in test_matrix_list])
nn_prediction_list = np.array([prediction(test_matrix, forest, under_mlp,
↳over_mlp)[0] for test_matrix in test_matrix_list])
actual_list = np.array([calculate_year_avg_pepm(claims.query('Client ==
↳@client'), 5, 1)[0] for client in client_list2])
py5_MSE = np.mean((prediction_list - actual_list) ** 2)
nn_py5_MSE = np.mean((nn_prediction_list - actual_list) ** 2)
original_py5_MSE = np.mean((model_predictions - actual_list) ** 2)
predictions = pd.DataFrame(data={'Predicted Values': prediction_list,
                                'Neural Network Predicted Values':
↳nn_prediction_list,
                                'Original Model Predicted Values':
↳model_predictions,
                                'Actual Values': actual_list,
                                'Client': client_list2})
display(predictions)
print(f'Combined MSE: {py5_MSE}')
print(f'Neural Network MSE: {nn_py5_MSE}')
print(f'Original MSE: {original_py5_MSE}')

```

	Predicted Values	Neural Network Predicted Values \
0	1116.106185	1126.906321
1	699.960189	711.200693
2	793.487099	804.697338
3	951.975490	1220.162770
4	647.469389	637.713426
5	831.170872	862.284041
6	1739.009898	1742.016076
7	934.464843	956.027727
8	802.947346	842.241710
9	668.376303	703.459771
10	687.176953	751.632823
11	1091.786803	1209.048180
12	817.325668	1068.409233

	Original Model Predicted Values	Actual Values \
0	1129.405610	1246.050814
1	634.706256	877.423241
2	788.990979	967.043199
3	1015.421540	1251.922508
4	603.845186	849.407154
5	872.817391	1035.917435
6	1770.223778	2213.431409
7	839.373682	843.197177
8	775.961080	1004.099845
9	622.383513	758.247882

10	635.381635	873.061295
11	1147.622884	1335.047505
12	874.037290	1029.380419

	Client
0	Antelope Valley
1	Avanti/Pipeline
2	Beverly Hospital
3	CHLA
4	Dameron
5	Enloe
6	Fairchild
7	Henry Mayo
8	Huntington Memorial Hospital
9	Prime Healthcare
10	Prospect Medical Holdings
11	Tahoe Forest
12	Torrance Health

Combined MSE: 51677.65017191957

Neural Network MSE: 33881.529993994096

Original MSE: 48443.10678834985

```
[24]: # Plan Year 6 Historical Predictions
client_list3 = claims.query('Client != "Epic" and Client != "Marshall" and
    ↳ Client != "CHA Hollywood Pres" and Client != "Northern Inyo" and Client !=
    ↳ "Salinas Valley" and Client != "Beverly Hospital" and Client !=
    ↳ "CHLA"')['Client'].unique()
first_three_list = [claims.query('Client == @client and (`Plan Year` == 3 or
    ↳ `Plan Year` == 4 or `Plan Year` == 5)') for client in client_list3]
test_matrix_list = [create_test_matrix(client_data, 3) for client_data in
    ↳ first_three_list]
model_predictions = np.array([model.predict(test_matrix)[0] for test_matrix in
    ↳ test_matrix_list])
prediction_list = np.array([prediction(test_matrix, forest, under_model,
    ↳ over_model)[0] for test_matrix in test_matrix_list])
nn_prediction_list = np.array([prediction(test_matrix, forest, under_mlp,
    ↳ over_mlp)[0] for test_matrix in test_matrix_list])
actual_list = np.array([calculate_year_avg_pepm(claims.query('Client ==
    ↳ @client'), 6, 1)[0] for client in client_list3])
py6_MSE = np.mean((prediction_list - actual_list) ** 2)
nn_py6_MSE = np.mean((nn_prediction_list - actual_list) ** 2)
original_py6_MSE = np.mean((model_predictions - actual_list) ** 2)
predictions = pd.DataFrame(data={'Predicted Values': prediction_list,
    ↳ 'Neural Network Predicted Values':
    ↳ nn_prediction_list,
```

```

                                'Original Model Predicted Values':
↪model_predictions,
                                'Actual Values': actual_list,
                                'Client': client_list3})
display(predictions)
print(f'Predicted Values MSE: {py6_MSE}')
print(f'Neural Network MSE: {nn_py6_MSE}')
print(f'Original MSE: {original_py6_MSE}')

```

	Predicted Values	Neural Network Predicted Values	\
0	1107.328750	1066.996501	
1	785.353241	802.381080	
2	701.186241	751.023962	
3	841.115553	727.285746	
4	1859.069131	1481.541633	
5	901.514849	966.672722	
6	857.826506	897.024992	
7	698.854908	750.386162	
8	790.677532	856.122491	
9	1074.938821	1443.929817	
10	817.285815	822.111590	

	Original Model Predicted Values	Actual Values	\
0	1177.151654	1315.931645	
1	719.439425	815.387330	
2	685.537568	747.456828	
3	882.410076	1137.423613	
4	1929.553526	2452.418010	
5	887.801248	931.659229	
6	821.077482	896.678063	
7	661.099666	842.884283	
8	747.306901	801.260753	
9	1101.222165	1398.492114	
10	866.422468	945.892936	

	Client
0	Antelope Valley
1	Avanti/Pipeline
2	Dameron
3	Enloe
4	Fairchild
5	Henry Mayo
6	Huntington Memorial Hospital
7	Prime Healthcare
8	Prospect Medical Holdings
9	Tahoe Forest
10	Torrance Health

Predicted Values MSE: 57356.42349559297
 Neural Network MSE: 109376.6170386512
 Original MSE: 46272.69708721688

```
[25]: # Plan Year 7 Historical Predictions
client_list4 = claims.query('Client != "Epic" and Client != "Marshall" and
    ↳ Client != "CHA Hollywood Pres" and Client != "Northern Inyo" and Client !=
    ↳ "Salinas Valley" and Client != "Beverly Hospital" and Client !=
    ↳ "CHLA")['Client'].unique()
first_three_list = [claims.query('Client == @client and (`Plan Year` == 4 or
    ↳ `Plan Year` == 5 or `Plan Year` == 6)') for client in client_list4]
test_matrix_list = [create_test_matrix(client_data, 4) for client_data in
    ↳ first_three_list]
model_predictions = np.array([model.predict(test_matrix)[0] for test_matrix in
    ↳ test_matrix_list])
prediction_list = np.array([prediction(test_matrix, forest, under_model,
    ↳ over_model)[0] for test_matrix in test_matrix_list])
nn_prediction_list = np.array([prediction(test_matrix, forest, under_mlp,
    ↳ over_mlp)[0] for test_matrix in test_matrix_list])
actual_list = np.array([calculate_year_avg_pepm(claims.query('Client ==
    ↳ @client'), 7, 1)[0] for client in client_list4])
py7_MSE = np.mean((prediction_list - actual_list) ** 2)
nn_py7_MSE = np.mean((nn_prediction_list - actual_list) ** 2)
original_py7_MSE = np.mean((model_predictions - actual_list) ** 2)
predictions = pd.DataFrame(data={'Predicted Values': prediction_list,
    'Neural Network Predicted Values':
    ↳ nn_prediction_list,
    'Original Model Predicted Values':
    ↳ model_predictions,
    'Actual Values': actual_list,
    'Client': client_list4})
display(predictions)
print(f'Predicted Values MSE: {py7_MSE}')
print(f'Neural Network MSE: {nn_py7_MSE}')
print(f'Original MSE: {original_py7_MSE}')
```

	Predicted Values	Neural Network Predicted Values \
0	1174.460883	1171.525945
1	811.675016	864.180295
2	701.704393	729.604004
3	926.140107	785.764096
4	2082.057446	1920.942340
5	905.490862	870.844548
6	834.271025	918.811524
7	777.896395	805.896901
8	813.268855	868.701444
9	1294.042317	1644.413975

10 849.228740 1050.155841

	Original	Model Predicted Values	Actual Values \
0		1212.078149	1319.483269
1		757.615902	712.878027
2		692.982914	918.272041
3		944.708423	1223.753275
4		2124.057917	2336.835668
5		832.379647	926.853781
6		827.395670	1035.579100
7		723.595852	816.691493
8		771.426865	895.584091
9		1241.425724	1644.965060
10		897.580673	961.426961

	Client
0	Antelope Valley
1	Avanti/Pipeline
2	Dameron
3	Enloe
4	Fairchild
5	Henry Mayo
6	Huntington Memorial Hospital
7	Prime Healthcare
8	Prospect Medical Holdings
9	Tahoe Forest
10	Torrance Health

Predicted Values MSE: 37834.25546277725

Neural Network MSE: 42787.74315066527

Original MSE: 39154.61740762532

```
[26]: # Plan Year 8 Historical Predictions
client_list5 = claims.query('Client != "Epic" and Client != "Marshall" and
↳ Client != "CHA Hollywood Pres" and Client != "Northern Inyo" and Client !=
↳ "Salinas Valley" and Client != "Beverly Hospital" and Client != "CHLA" and
↳ Client != "Avanti/Pipeline" and Client != "Enloe" and Client != "Huntington
↳ Memorial Hospital" and Client != "Prospect Medical Holdings" and Client !=
↳ "Tahoe Forest")['Client'].unique()
first_three_list = [claims.query('Client == @client and (`Plan Year` == 5 or
↳ `Plan Year` == 6 or `Plan Year` == 7)') for client in client_list5]
test_matrix_list = [create_test_matrix(client_data, 5) for client_data in
↳ first_three_list]
model_predictions = np.array([model.predict(test_matrix)[0] for test_matrix in
↳ test_matrix_list])
prediction_list = np.array([prediction(test_matrix, forest, under_model,
↳ over_model)[0] for test_matrix in test_matrix_list])
```

```

nn_prediction_list = np.array([prediction(test_matrix, forest, under_mlp,
↳over_mlp)[0] for test_matrix in test_matrix_list])
actual_list = np.array([calculate_year_avg_pepm(claims.query('Client ==
↳@client'), 8, 1)[0] for client in client_list5])
py8_MSE = np.mean((prediction_list - actual_list) ** 2)
nn_py8_MSE = np.mean((nn_prediction_list - actual_list) ** 2)
original_py8_MSE = np.mean((model_predictions - actual_list) ** 2)
predictions = pd.DataFrame(data={'Predicted Values': prediction_list,
                                'Neural Network Predicted Values':
↳nn_prediction_list,
                                'Original Model Predicted Values':
↳model_predictions,
                                'Actual Values': actual_list,
                                'Client': client_list5})
display(predictions)
print(f'Predicted Values MSE: {py8_MSE}')
print(f'Neural Network MSE: {nn_py8_MSE}')
print(f'Original MSE: {original_py8_MSE}')

```

	Predicted Values	Neural Network Predicted Values \
0	1136.863873	1110.901211
1	837.705299	862.688675
2	2031.850184	1904.721175
3	869.518138	899.426386
4	786.981548	831.490667
5	858.617069	1036.698482

	Original Model Predicted Values	Actual Values	Client
0	1210.216238	1561.734835	Antelope Valley
1	814.442635	1022.715639	Dameron
2	2117.359881	2570.400823	Fairchild
3	835.245193	824.657295	Henry Mayo
4	748.439797	920.782757	Prime Healthcare
5	907.070020	1066.047820	Torrance Health

Predicted Values MSE: 94620.62122743168
 Neural Network MSE: 114402.30972695122
 Original MSE: 71212.86739033448

```

[27]: # Plan Year 9 Historical Predictions
client_list6 = claims.query('Client != "Epic" and Client != "Marshall" and
↳Client != "CHA Hollywood Pres" and Client != "Northern Inyo" and Client !=
↳"Salinas Valley" and Client != "Beverly Hospital" and Client != "CHLA" and
↳Client != "Avanti/Pipeline" and Client != "Enloe" and Client != "Huntington
↳Memorial Hospital" and Client != "Prospect Medical Holdings" and Client !=
↳"Tahoe Forest" and Client != "Fairchild" and Client != "Henry
↳Mayo")['Client'].unique()

```

```

first_three_list = [claims.query('Client == @client and (`Plan Year` == 6 or
↳ `Plan Year` == 7 or `Plan Year` == 8)) for client in client_list6]
test_matrix_list = [create_test_matrix(client_data, 6) for client_data in
↳ first_three_list]
model_predictions = np.array([model.predict(test_matrix)[0] for test_matrix in
↳ test_matrix_list])
prediction_list = np.array([prediction(test_matrix, forest, under_model,
↳ over_model)[0] for test_matrix in test_matrix_list])
nn_prediction_list = np.array([prediction(test_matrix, forest, under_mlp,
↳ over_mlp)[0] for test_matrix in test_matrix_list])
actual_list = np.array([calculate_year_avg_pepm(claims.query('Client ==
↳ @client'), 9, 1)[0] for client in client_list6])
py9_MSE = np.mean((prediction_list - actual_list) ** 2)
nn_py9_MSE = np.mean((nn_prediction_list - actual_list) ** 2)
original_py9_MSE = np.mean((model_predictions - actual_list) ** 2)
predictions = pd.DataFrame(data={'Predicted Values': prediction_list,
                                'Neural Network Predicted Values':
↳ nn_prediction_list,
                                'Original Model Predicted Values':
↳ model_predictions,
                                'Actual Values': actual_list,
                                'Client': client_list6})
display(predictions)
print(f'Predicted Values MSE: {py9_MSE}')
print(f'Neural Network MSE: {nn_py9_MSE}')
print(f'Original MSE: {original_py9_MSE}')

```

	Predicted Values	Neural Network Predicted Values	\
0	1272.976154	1367.846237	
1	888.549705	905.572455	
2	864.437722	920.284777	
3	903.969212	994.246758	

	Original Model Predicted Values	Actual Values	Client
0	1302.446083	1599.365944	Antelope Valley
1	843.541169	1157.420117	Dameron
2	796.584788	914.276404	Prime Healthcare
3	929.375044	1049.281601	Torrance Health

Predicted Values MSE: 50605.29469734833

Neural Network MSE: 30023.388540240783

Original MSE: 53727.57426073433

[28]: *# Plan Year 10 Historical Predictions*

```

first_three_list = [claims.query('Client == @client and (`Plan Year` == 7 or
↳ `Plan Year` == 8 or `Plan Year` == 9)) for client in client_list6]

```



```

test_matrix_list = [create_test_matrix(client_data, 7) for client_data in
    ↪first_three_list]
model_predictions = np.array([model.predict(test_matrix)[0] for test_matrix in
    ↪test_matrix_list])
prediction_list = np.array([prediction(test_matrix, forest, under_model,
    ↪over_model)[0] for test_matrix in test_matrix_list])
nn_prediction_list = np.array([prediction(test_matrix, forest, under_mlp,
    ↪over_mlp)[0] for test_matrix in test_matrix_list])
actual_list = np.array([calculate_year_avg_pepm(claims.query('Client ==
    ↪@client'), 10, 1)[0] for client in client_list6])
py10_MSE = np.mean((prediction_list - actual_list) ** 2)
nn_py10_MSE = np.mean((nn_prediction_list - actual_list) ** 2)
original_py10_MSE = np.mean((model_predictions - actual_list) ** 2)
predictions = pd.DataFrame(data={'Predicted Values': prediction_list,
    'Neural Network Predicted Values':
    ↪nn_prediction_list,
    'Original Model Predicted Values':
    ↪model_predictions,
    'Actual Values': actual_list,
    'Client': client_list6})
display(predictions)
print(f'Predicted Values MSE: {py10_MSE}')
print(f'Neural Network MSE: {nn_py10_MSE}')
print(f'Original MSE: {original_py10_MSE}')

```

	Predicted Values	Neural Network Predicted Values	\
0	1345.066825	1382.850943	
1	998.784340	1060.122971	
2	861.021503	940.687134	
3	878.576641	956.190212	

	Original Model Predicted Values	Actual Values	Client
0	1392.792525	1266.928222	Antelope Valley
1	941.180078	1038.074140	Dameron
2	823.507042	920.893849	Prime Healthcare
3	939.581742	1062.023598	Torrance Health

Predicted Values MSE: 11221.70336779696
 Neural Network MSE: 6379.177002677565
 Original MSE: 12426.620027321846

```

[29]: # Plan Year 11 Historical Predictions
first_three_list = [claims.query('Client == @client and (`Plan Year` == 8 or
    ↪`Plan Year` == 9 or `Plan Year` == 10)') for client in client_list6]
test_matrix_list = [create_test_matrix(client_data, 8) for client_data in
    ↪first_three_list]

```

```

model_predictions = np.array([model.predict(test_matrix)[0] for test_matrix in
    ↪test_matrix_list])
prediction_list = np.array([prediction(test_matrix, forest, under_model,
    ↪over_model)[0] for test_matrix in test_matrix_list])
nn_prediction_list = np.array([prediction(test_matrix, forest, under_mlp,
    ↪over_mlp)[0] for test_matrix in test_matrix_list])
actual_list = np.array([calculate_year_avg_pepm(claims.query('Client ==
    ↪@client'), 11, 1)[0] for client in client_list6])
py11_MSE = np.mean((prediction_list - actual_list) ** 2)
nn_py11_MSE = np.mean((nn_prediction_list - actual_list) ** 2)
original_py11_MSE = np.mean((model_predictions - actual_list) ** 2)
predictions = pd.DataFrame(data={'Predicted Values': prediction_list,
    'Neural Network Predicted Values':
    ↪nn_prediction_list,
    'Original Model Predicted Values':
    ↪model_predictions,
    'Actual Values': actual_list,
    'Client': client_list6})

display(predictions)
print(f'Predicted Values MSE: {py11_MSE}')
print(f'Neural Network MSE: {nn_py11_MSE}')
print(f'Original MSE: {original_py11_MSE}')

```

	Predicted Values	Neural Network Predicted Values	\
0	1354.912931	2039.597381	
1	970.057035	1154.312004	
2	921.571788	981.388061	
3	983.679294	1214.067199	

	Original Model Predicted Values	Actual Values	Client
0	1411.401673	1266.928222	Antelope Valley
1	989.224653	1444.641882	Dameron
2	871.780133	951.959321	Prime Healthcare
3	998.061455	1434.575063	Torrance Health

Predicted Values MSE: 109300.62075722024

Neural Network MSE: 182699.7092062607

Original MSE: 106312.56578824003

1.5 Conclusion

Comparing performance on historical data, it becomes immediately obvious that this model falls short of its original goal to predict yearly healthcare costs from 36 months of claims experience data. Although there was not enough time to conduct a full analysis on this model's specific shortcomings, there are some indications within the report itself. The first, most obvious, is the shape of the residual plots (both of the original Ridge Regression model and the subsequent under and overestimation models). Each featured some sort of positive correlation which indicates the features used to train the model do not account for all of the observed variance in the data. A

more careful analysis of data available in Mede to create a more complex model would likely result in a better model. Second is a lack of data. Although the original `claims` data included some 1500 rows of data, the granularity was monthly. Given 36 months of data were required to predict, this reduced the number of unique datapoints to about 40. While the statistical methods applied are sound, having a more representative data sample is always better.

Overall, although the final model's utility is perhaps less than anticipated, as noted at the beginning of the report, this analysis is just a starting point. During lunch with Mitch at some point during my time here, he brought up the need for more advanced analytics in the insurance industry. It was a sentiment echoed by Ju at the June Healthcare meeting where I was first introduced to the rest of the team and one I wholeheartedly agree with. Although this model is not one that can necessarily be used, I have no doubt that with more work it could be refined into a model that outperforms current models/procedures for calculating renewal projections across the board. It is my hope that this analysis/report demonstrates the value that Machine Learning driven algorithms have in this industry and serves as a starting point for bringing more advanced analytics to Keenan. As I expressed in my final reflections email, for better or worse the world is continually finding new applications for more powerful predictive modelling and Keenan/AP must keep up to continue to stay ahead of its competitors.